

PASS4SURES.COM

A Composite Solution With Just One Click

Microsoft

70-466 PRACTICE EXAM

Microsoft Implementing Data Models and Reports with Microsoft SQL Server

2012 Exam

Product Questions: 188

Version: 17.0

Case Study: 1, Tailspin Toys

Tailspin Toys is a multinational company that manufactures toys. Tailspin Toys has offices in five regions worldwide. The company sells toys at various retail stores. The company also sells toys directly to consumers through a web site.

The company has the following departments:

- Sales
- Distribution
- Manufacturing

Each department has an office in each region.

The fiscal calendar of Tailspin Toys runs from June to May.

The network contains a server farm that has Microsoft SharePoint Server 2013 installed.

Existing Environment

Current Database Environment

Each department uses SharePoint team sites for internal collaboration.

All manufacturing information is stored in a relational database named Manufacturing. All sales information is stored in a relational database named Sales.

Tailspin Toys deploys SQL Server Analysis Services (SSAS) and configures SSAS to use tabular models. SSAS will be used for all sales reports.

Tailspin Toys deploys a SQL Server Reporting Services (SSRS) instance in SharePoint mode.

Sales Database

A database named Sales contains two tables named FactSales and DimProduct. FactSales contains the following columns:

- SalesID
- Total Due
- OrderDate

DimProduct contains the following columns:

- ProductID
- ProductName
- ProductCategory
- ProductSubcategory

The Sales database contains information about the products. Most of the products have a category and a subcategory. Certain products only have a category.

A sample from DimProduct is shown in the following table.

ProductID	ProductName	ProductCategory	ProductSubcategory
1	Balsa Wood Flyer	Plane	Classic
2	Radio Controlled Flyer	Plane	Radio Controlled
3	Plastic Model	Model	Model

Requirements

Security Requirements

Tailspin Toys identifies the following security requirement:

- Sales department users must be allowed to view the sales transactions from their region only.
- Sales department users must be able to view the contents of the manufacturing reports.
- Manufacturing department users must be able to create new manufacturing reports.
- Third-party and custom solutions must NOT be deployed to the reporting server.
- Sales department users must NOT be able to create new manufacturing reports.

Planned Reporting Implementation

The manufacturing department plans to use the SSRS instance for its reports. The manufacturing department also plans to make its reports accessible from SharePoint. All manufacturing reports will use an existing database named Manufacturing.

Reporting Requirements

Tailspin Toys identifies the following reporting requirements:

- All reports must contain the company logo and a header that contains the date and the time that the report was executed.
- All reports must be created by using the SQL Server Data Tools.

Manufacturing report

You plan to create a report named ManufacturingIssues.rdl. The report has the following requirements:

- Manufacturing department managers must be able to view product issues by product type, manufacturing plant location, and error type.

- The manufacturing department managers must be able to change views by choosing options from drop-down lists.

Sales reports

You plan to create a sales report named RegionalSales.rdl. The report has the following requirements:

- Users must be able to view the report by using a web browser. By default, subcategories and product details must be hidden when using the browser.
- Users must be able to subscribe to receive the report by email. The report must be sent by email as a PDF attachment.

You plan to create a quarterly sales report named QuarterSales.rdl. The report must display sales data by fiscal quarter.

Technical Requirements

Tailspin Toys identifies the following technical requirements:

- Products in the DimProduct table that do NOT have a subcategory must use the category value as the subcategory value.
- SSRS must NOT connect to databases more frequently than once every 30 minutes.
- Sales department users must be able to use Microsoft Excel to browse tabular data.

Question: 1

You need to recommend a solution for the sales department that meets the security requirements. What should you recommend?

- A. Create one role for all of the sales department users. Add a DAX filter that reads the current user name and retrieves the user's region.
- B. Create one role for each region. Configure each role to have read access to a specific region. Add the sales department users to their corresponding role.
- C. Create a table for each region. Create a role for each region. Grant each role read access to its corresponding table.
- D. Create one role for all of the sales department users. Configure the role to have read access to the sales transactions. Ensure that all of the reports that access the sales transaction data restrict read access to the data from the corresponding sales department region only.

Answer: C

Explanation:

Scenario: Tailspin Toys identifies the following security requirement:

- Sales department users must be allowed to view the sales transactions from their region only.
- Sales department users must be able to view the contents of the manufacturing reports.
- Sales department users must NOT be able to create new manufacturing reports.

Question: 2

You need to configure the dataset for the ManufacturingIssues report. The solution must meet the technical

requirements and the reporting requirements.

What should you do?

- A. Configure the dataset to use a stored procedure. Add the necessary parameters to the stored procedure.
- B. Add a query to retrieve the necessary data from the database. Configure the dataset to use query parameters.
- C. Add a query to retrieve the necessary data from the database. Configure the dataset to use filter parameters.
- D. Configure the dataset to use a table. Ensure that the database has a table that contains the necessary information.

Answer: B

Question: 3

You need to ensure that all reports meet the reporting requirements.

What is the best way to achieve the goal? More than one answer choice may achieve the goal. Select the BEST answer.

- A. Create a report part. Publish the report part to a server that has SSRS installed. Add the report part to each new report that is created.
- B. Create a report part. Publish the report part to a SharePoint site. Add the report part to each new report that is created.
- C. Create a report. Copy the report to source code control. Create each new report by using the report template in source code control.
- D. Create a report. Copy the report to the PrivateAssemblies\ProjectItems\ReportProject folder in the Visual Studio directory. Create each new report by using the locally stored report

Answer: D

Question: 4

You need to configure a hierarchy for DimProduct that meets the technical requirements.

What should you do?

- A. Set ProductName as the parent of ProductSubCategory and set ProductSubcategory as the parent of ProductCategory. For ProductSubcategory, click Hide if Name Equals Parent.
- B. Set ProductCategory as the parent of ProductSubCategory and set ProductSubcategory as the parent of ProductName. For ProductSubcategory, click Hide if Name Equals Parent.
- C. Set ProductName as the parent of ProductSubcategory and set ProductSubCategory as the parent of ProductCategory. For ProductCategory, click Hide if Name Equals Parent.
- D. Set ProductCategory as the parent of ProductSubcategory and set ProductSubCategory as the parent of ProductName. For ProductCategory, click Hide if Name Equals Parent.

Answer: B

Question: 5

You need to recommend a solution to meet the requirements for the ManufacturingIssues.rdl report.

What is the best solution that you should include in the recommendation? More than one answer choice may achieve the goal. Choose the BEST answer.

- A. Add a dataset to the report that uses an ad hoc SQL statement. Configure the dataset to include the parameters required for the different views. Add a dataset for each parameter created. Configure each parameter to use the values in the dataset.
- B. Add a dataset to the report that uses an ad hoc SQL statement. Configure the dataset to include the parameters required for the different views. Update each parameter to use a set of values from Report Designer.
- C. Add a dataset to the report that uses an ad hoc SQL statement. Configure the dataset to include the parameters required for the different views. Use the default display for the parameters.
- D. Add a dataset to the report that uses a stored procedure. Configure the dataset to include the parameters required for the different views. Update each parameter to use a set of values from Report Designer.

Answer: C

Question: 6

HOTSPOT

You create a new SharePoint site to store reports for the manufacturing department.

You need to recommend a solution to meet the security requirements for the sales department users and the manufacturing department users.

What should you recommend? To answer, select the appropriate group for the sales department users and the manufacturing department users in the answer area.

Answer Area

Sales department users	Manufacturing department users
<input type="text"/>	<input type="text"/>

Answer Area

Sales department users	Manufacturing department users
<input type="text"/> Owners Readers Members Restricted readers	<input type="text"/> Visitors Readers Members Restricted readers

Answer:

Answer Area

Sales department users	Manufacturing department users
<p>Owners Readers Members Restricted readers</p>	<p>Visitors Readers Members Restricted readers</p>

Question: 7

You need to modify the environment before you create the QuarterSales report.

What should you do?

- A. Add a date table to the model that contains columns for the fiscal and calendar quarters.
- B. Add a date table to the model that contains measures for the fiscal and calendar quarters.
- C. Configure a time dimension by using the Time Intelligence Wizard.
- D. Configure SSAS to use a server time dimension.

Answer: C**Question: 8**

After you deploy the RegionalSales report, users report that they cannot see product data when they receive the reports by email.

You need to ensure that the sales department managers can see all of the data.

In the report, you update the Hidden property of each group.

What should you do next?

- A. When the report is initially run, select Show or hide based on an expression. Set the expression to =
(Globals!RenderFormat.IsInteractive)
- B. When the report is initially run, select Show.
- C. When the report is initially run, select Show or hide based on an expression. Set the expression to =NOT
(Globals!RenderFormat.IsInteractive).
- D. When the report is initially run, select Hide.

Answer: A**Question: 9**

After you deploy the RegionalSales report, you attempt to configure the subscriptions.

You discover that the subscription creation screen does not display the option to deliver the report by email.

You need to ensure that subscriptions can be delivered by using email.

What should you do?

- A. Modify the Rsmgrpolicy.config file.
- B. From Central Administration, modify the SMTP settings of the SharePoint Server server farm.
- C. Modify the Rssrvrpolcy.config file.
- D. From Central Administration, modify the properties of the Reporting Service Application.

Answer: B

**Case Study: 2,
Contoso, Ltd****Background**

You are the business intelligence (BI) solutions architect for Contoso Ltd, a multinational sales company with offices in London, Madrid, Paris, Brisbane, Tokyo, and New York. Contoso sells office consumable products such as pens, printer ink, and paper.

You produce solutions by using SQL Server 2012 Business Intelligence Edition and Microsoft SharePoint Server 2010 Enterprise Edition with SP1.

Technical Background

Contoso's products are categorized by using four levels while some use only two or three levels. Products are categorized as shown in the following table.

Product Type	Product Category	Product Sub Category	Product Sub Section
Papers	Copy Paper		
	Note	Sticky Notes	
		"Sign Here" Notes	
Tapes and Glue	Adhesive Glue		
	Tape	Masking Tape	
		Sticky Tape	
Writing	Pens	Ball Pens	
		Pencils	
		WhiteBoard Markers	Permanent Markers
	Corrections		Removable Markers
		Correction Tape	
		Correction Fluid	
	Erasers		

Contoso sells products through mobile sales staff, direct marketing, and its website. Sales personnel are located in various regions around the world, and each region has a sales manager who is paid a quarterly bonus based on the total sales in the region during the quarter. Regions are categorized as shown in the following table.

Region	Country	State
Oceania	Australia	Queensland
		New South Wales
	New Zealand	Canterbury
Europe		Marlborough
	Great Britain	Cornwall
		Aberdeen
		Cardiff
	Germany	Baden-Wurttemberg
		Saxony

SQL Server Analysis Services (SSAS) is used to host a multidimensional database. The database contains a single cube named Sales and three database dimensions named Products, Regions, and Date. A single measure named Sales Total has been defined in the cube. The data source for the database is a SQL Server data warehouse.

The Products dimension contains a single user-defined hierarchy named Products. To prevent the display of empty members when users browse the Products dimension, the Extract, Transform, and Load (ETL) process populates all missing values as shown in the following diagram.

Product Type	Product Category	Product Sub Category	Product Sub Section
Papers	Copy Paper	Copy Paper	Copy Paper
Papers	Note Papers	Sticky Notes	Sticky Notes

The structure of the Products hierarchy is shown in the following diagram.

Products
Product Type
Product Category
Product Sub Category
Product Sub Section

The Regions dimension contains a single user-defined hierarchy named Sales Regions. The dimension is based on a single dimension table in the data warehouse and the attribute relationships have not been modified since the dimension was created by using the Dimension wizard. The structure of the Sales Regions hierarchy is shown in the following diagram.

Sales Regions
Region
Country
State

The Date dimension contains a single user-defined hierarchy named Calendar. The structure of the Calendar hierarchy is shown in the following diagram.

Calendar
Year
Quarter
Month
Date

A role named UserRegions has been created in the SSAS database that will be used to filter members in the Regions dimension based on the authenticated user.

Administrative staff from around the world will produce sales reports with Microsoft Excel 2010 based on the Sales cube.

Developers will produce reports with SQL Server Reporting Services (SSRS) based on the Sales cube and the reports will be delivered to users through report subscriptions and a web browser-All users log on to an Active Directory Domain Services (AD DS) domain named contoso.com.

All client computers and servers are joined to the contoso.com domain.

Business Requirements

The BI system must meet the following reporting requirements:

- Display all sales figures in euro currency, regardless of the client's reporting location
- Include a new measure named AD Sales that calculates average daily sales for a selected month
- Support near real-time reporting while maintaining good performance for multidimensional queries
- Support reports that show currency exchange rates
- Deliver executive reports that are parameterized and rendered from report snapshots

In addition, cube objects must use terms familiar to users from around the world. For example, in the SalesRegions hierarchy, users from Great Britain must see the State level presented as County when browsing the Sales cube.

The Sales cube must support a new measure group named Sales Planning. The measure group must consist of a single measure named Sales Plan that enables the management team to use Excel 2010 to enter sales plans for future monitoring.

The BI system must meet the following technical requirements:

Architecture requirements

- The system must use separate servers for each of the following components:
 - SQL Server Database Engine
 - SQL Server Integration Services
 - SQL Server Analysis Services in multidimensional mode
 - SharePoint Server with the Reporting Services Add-in
- All servers must be installed using U.S. regional settings.
- The system must source currency exchange rate data from a database hosted in Microsoft Azure SQL Database.

Security requirements

- When possible, the system must use Windows authentication for all database connections.
- The system must prevent users from querying data from outside of their region.
- The system must allow certain users to query data from multiple regions.

Development requirements

- When browsing the Products hierarchy, repeating values for different levels of a given drill-path must be avoided. For example, Papers -> Copy Paper -> Copy Paper -> Copy Paper should appear simply as Papers -> Copy Paper.
- The system must support report snapshots. The default maximum number of retained snapshots must not exceed five.

Question: 1

You need to develop an SSRS report that retrieves currency exchange rate data.
How should you configure the data source for the report?

- A. Use the Microsoft Azure SQL Database data source type and then set a username and password for the credentials.
- B. Use the SQL Server data source type and then set Windows authentication for the credentials.
- C. Use the Microsoft Azure SQL Database data source type and then set Windows authentication for the credentials.
- D. Use the SQL Server data source type and then set a username and password for the credentials.

Answer: A

Question: 2

You need to create the AD Sales measure.
Which aggregation function should you use?

- A. Sum
- B. Average
- C. ByAccount
- D. AverageOfChildren

Answer: D

Question: 3

You need to meet the browsing requirements for the Products hierarchy.
Which property should you modify?

- A. DefaultMember
- B. AttributeHierarchyDisplayFolder
- C. HideMemberIf
- D. RootMemberIf

Answer: C

Question: 4

You need to configure the partition storage settings to support the reporting requirements.
Which partition storage setting should you use?

- A. High-latency MOLAP
- B. In-Memory
- C. Low-latency MOLAP
- D. DirectQuery
- E. LazyAggregations
- F. Regular

Answer: C

Question: 5

You need to configure per-user security authentication for reporting against the Sales cube.
What should you do? (Each correct answer presents part of the complete solution. Choose all that apply.)

- A. Create Service Principal Names (SPNs).
- B. Enable forms-based authentication.
- C. Configure account delegation.
- D. Enable mixed-mode authentication.

Answer: A, D

Question: 6

You need to configure SSRS to meet the maximum number of snapshots requirement.
What should you do? (Each answer presents a complete solution. Choose all that apply.)

- A. In SharePoint Central Administration, set the System Snapshot Limit option to 5.
- B. In Reporting Services Configuration Manager, set the Limit number of snapshots option to 5.
- C. For each report, set the System Snapshot Limit option to 5.

D. Use PowerShell to set the System Snapshot Limit option to 5.

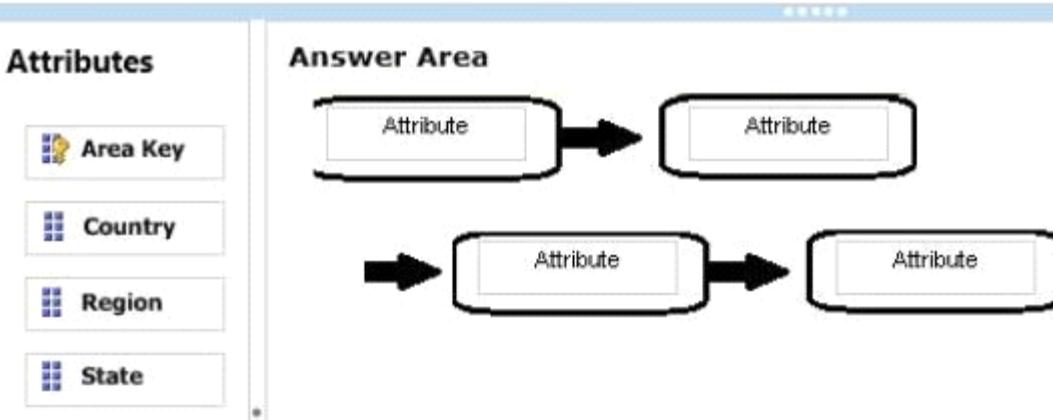
Answer: A, C

Question: 7 DRAG

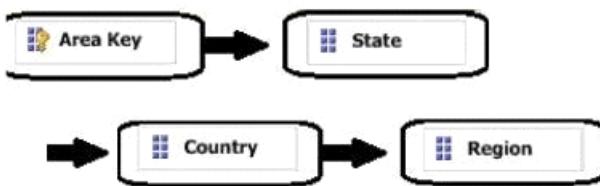
DROP

You need to optimize the Regions dimension.

How should you configure the attribute relationships? (To answer, drag the appropriate attribute from the list of attributes to the correct location in the hierarchy relationship. Use only attributes that apply.)



Answer:



Question: 8

You need to modify the Sales cube to support the planning requirements.

Which SSAS feature should you use?

- A. A KPI
- B. A translation
- C. A perspective
- D. A writeback partition

Answer: D

Question: 9

You need to modify the Sales Regions hierarchy to meet the reporting requirements.

Which SSAS feature should you use?

- A. Calculation
- B. Translation

- C. Perspective
- D. Action

Answer: B

Question: 10

You need to develop the executive reports.

What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

- A. Provide default values for all parameters.
- B. Set the data source to use Windows authentication.
- C. Remove default values from all parameters.
- D. Implement dataset filters to filter data.
- E. Set the data source to use stored Windows credentials.
- F. Implement dataset query parameters to filter data.

Answer: A, D, E

Question: 11

You need to configure the UserRegions role.

Which Multidimensional Expressions (MDX) function should you use?

- A. USERNAME()
- B. USERID()
- C. CUSTOMDATA()
- D. UNIQUENAME()
- E. LOOKUPVALUE()

Answer: A

Question: 12

You need to configure the format of the Sales Total measure.

Which value should you use for the FormatString property?

- A. \$#,##0.00;(\$#,##0.00)
- B. #,##0.00;-#,##0.00
- C. Currency
- D. A custom-entered value

Answer: C

Case Study: 3, Data Architect General Background

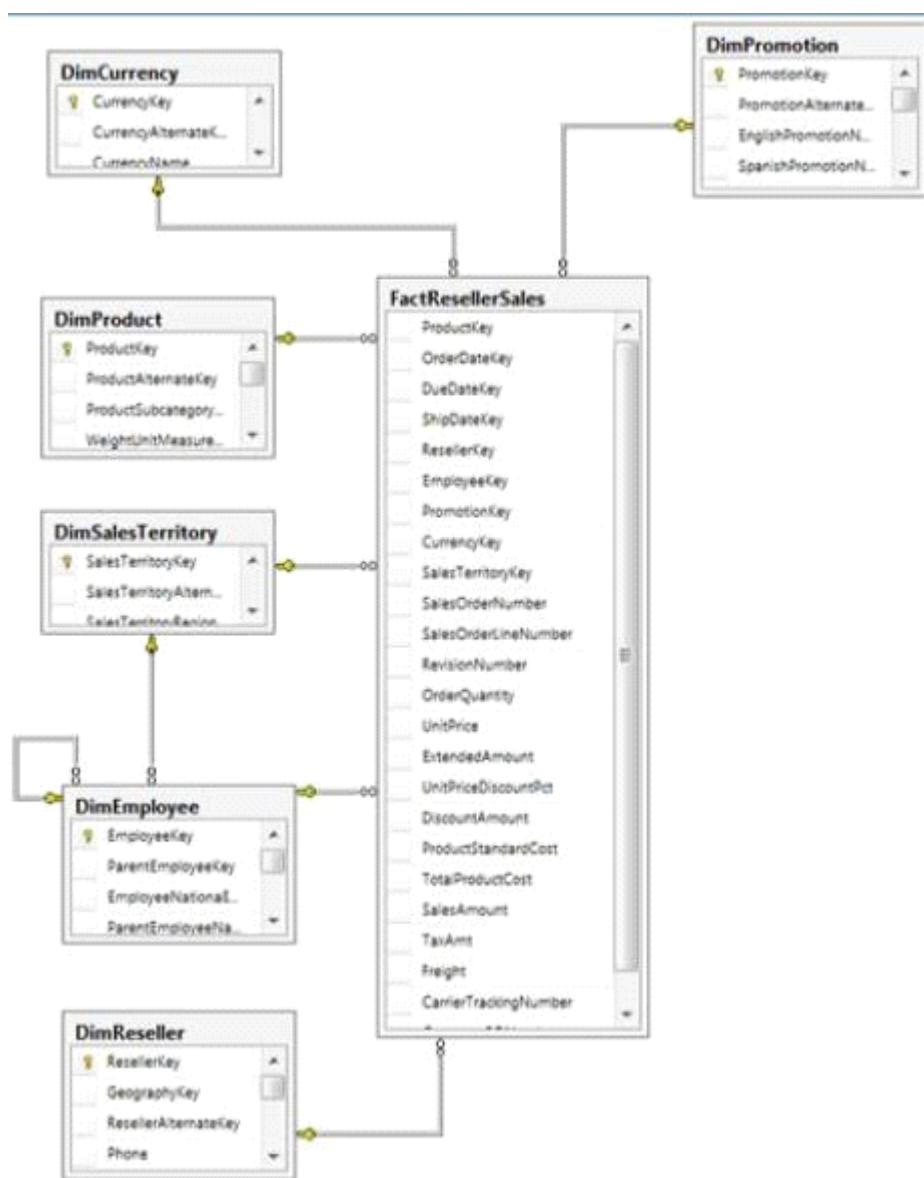
You are the data architect for a company that uses SQL Server 2012 Enterprise Edition. You design data modeling and reporting solutions that are based on a sales data warehouse.

Background

The solutions will be deployed on the following servers:

- ServerA runs SQL Server Database Engine, ServerA is the data warehouse server.
- ServerB runs SQL Server Database Engine, SQL Server Analysis Services (SSAS) in multidimensional mode, and SQL Server Integration Services (SSIS).
- ServerC runs SSAS in tabular mode, SQL Server Reporting Services (SSRS) running in SharePoint mode, and Microsoft SharePoint 2010 Enterprise Edition with SP1.

The data warehouse schema currently contains the tables shown in the exhibit. (Click the Exhibit button.)



Business Requirements

The reporting solution must address the requirements of the sales team, as follows:

- Team members must be able to view standard reports from SharePoint.
- Team members must be able to perform ad-hoc analysis by using Microsoft Power View and Excel.
- Team members can have standard reports delivered to them on a schedule of their choosing.

The standard reports

- Will use a sales territory hierarchy for organizing data by region.
- Will be accessible from SharePoint.

The Excel ad-hoc reports

- Will use the same data store as the standard reports.
- Will provide direct access to the data store for the sales team and a simplified view for the executive team.

Technical Requirements

The standard reports must be based on an SSAS cube. The schema of the data warehouse on ServerA must be able to support the ability to slice the fact data by the following dates:

- Order date (OrderDateKey)
- Due date (DueDateKey)
- Ship date (ShipDateKey)

Additions and modifications to the data warehouse schema must adhere to star schema design principles to minimize maintenance and complexity

The multidimensional and tabular models will be based on the data warehouse. The tabular and multidimensional models will be created by using SQL Server Data Tools (SSDT). The tabular project is named AdhocReports and the multidimensional project is named Standard Reports.

The cube design in the Standard Reports project must define two measures for the unique count of sales territories (SalesTerritoryKey) and products (ProductKey).

A deployment script that can be executed from a command-line utility must be created to deploy the StandardReports project to ServerB.

The tabular model in the AdhocReports project must meet the following requirements:

- A hierarchy must be created that consists of the SalesTerritoryCountry and SalesTerritoryRegion columns from the DimSalesTerritory table and the EmployeeName column from the DimEmployee table.
- A key performance indicator (KPI) must be created that compares the total quantity sold (OrderQuantity) to a threshold value of 1,000.
- A measure must be created to calculate day-over-day (DOD) sales by region based on order date.

SSRS on ServerC must be configured to meet the following requirements:

- It must use a single data source for the standard reports.
- It must allow users to create their own standard report subscriptions.

- The sales team members must be limited to only viewing and subscribing to reports in the Sales Reports library.

A week after the reporting solution was deployed to production, Marc, a salesperson, indicated that he has never received reports for which he created an SSRS subscription. In addition, Marc reports that he receives timeout errors when running some reports on demand.

Question: 1

You need to create the KPI in the AdhocReports project.
What should you do?

- Create a measure by using the SUM([OrderQuantity]) expression. Then use the CREATE KPI CURRENTCUBE statement to define the KPI and target value.
- Create a measure by using the SUM([OrderQuantity]) expression and create a KPI based on the measure. Then set the target value.
- Create a measure by using the COUNT([OrderQuantity]) expression and create a KPI based on the measure. Then set the target value.
- Create a KPI based on the OrderQuantity column and then set the target value.

Answer: A

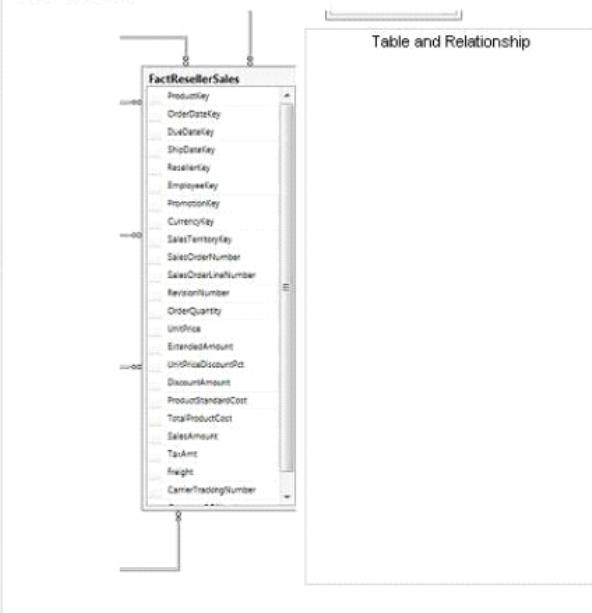
Question: 2

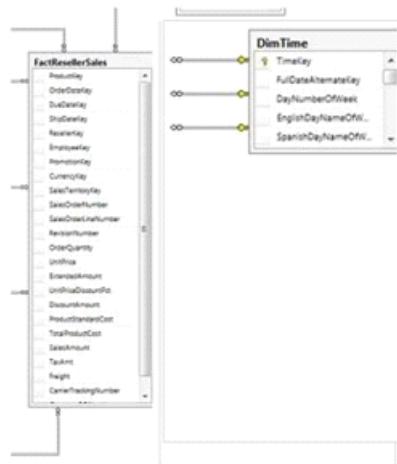
DRAG DROP

You need to complete the design of the data warehouse.

Which design should you use? (To answer, drag the appropriate tables and relationships to the correct location in the answer area.)

- Use only the tables and relationships that apply.)

Tables and Relationships**Answer Area****Answer:**



Question: 3

You need to create the sales territory and product measures.
Which aggregate function should you use for both measures?

- A. Count
- B. COUNT(DISTINCT column_name)
- C. DistinctCount
- D. Distinct

Answer: C

Question: 4

You need to ascertain why Marc did not receive his reports.
What should you do?

- A. Search the ReportServerService_{timestamp}.log file for errors.
- B. Search the registry for errors.
- C. Use SQL Server Management Studio to search the SQL Server logs for errors.
- D. Use the Windows Event Viewer to search the Application log for errors.

Answer: A

Question: 5

You need to create the hierarchy in the AdhocReports project.
What should you do?

- A. Multi-select all of the columns, right-click the columns, and then click the Create Hierarchy command.
- B. Use the RELATEDTABLE() function to consolidate the tables, multi-select the columns in the hierarchy,

- right-click the columns, and then click the Create Hierarchy command.
- C. Use the RELATED() function to consolidate the columns in the DimSalesTerritory table, multi-select the columns, right-click the columns, and then click the Create Hierarchy command.
- D. Use the RELATED() function to consolidate the columns in the DimEmployee table, multi-select the columns, right-click the columns, and then click the Create Hierarchy command.

Answer: D

Question: 6

- You need to create the KPI in the AdhocReports project in time for the next production release cycle. What should you do?
- A. Create a measure by using the COUNT([OrderQuantity]) expression and create a KPI based on the measure. Then set the target value. Check in the changes before the next release cycle.
- B. Create a KPI based on the OrderQuantity column and then set the target value. Check in the changes before the next release cycle.
- C. Create a measure by using the SUM([OrderQuantity]) expression and create a KPI based on the measure. Then set the target value. Check in the changes before the next release cycle.
- D. Create a measure by using the SUM([OrderQuantity]) expression. Then use the CREATE KPI CURRENTCUBE statement to define the KPI and target value. Check in the changes before the next release cycle.

Answer: D

Question: 7

- You need to identify the reports that produce the errors that Marc is receiving. What should you do?

- A. Write a query by using the Subscriptions table in the report server database.
- B. Write a query by using the Execution Log 3 view in the report server database.
- C. Use the Windows Event Viewer to search the Application log for errors.
- D. Search the ReportServerService_<timestamp>.log file for errors.

Answer: B

Question: 8

- You need to create the data source view for the StandardReports project. What should you do?

- A. Create a data source, connect it to the data warehouse, and then use the Data Source View wizard.

- B. Generate a relational schema from the dimensions and cubes by using the Schema Generation wizard.
- C. Create a new data source view and then use the Import from Table wizard.
- D. Execute the Import from Table wizard and then use the Data Source View wizard.

Answer: A

Question: 9

You need to create a measure for DOD sales.

What should you do? (Each correct answer presents part of the solution. Choose all that apply.)

- A. Use the Data Analysis Expressions (DAX) PARALLELPERIOD () function.
- B. Create a date dimension by using the Dimension wizard with a date template.
- C. Specify a date table by using a Mark the Date table.
- D. Use the Multidimensional Expressions (MDX) PARALLELPERIOD() function.

Answer: A, C

Question: 10

You need to develop the multidimensional project to meet the requirements of the Excel users.

What should you do?

- A. Create a separate cube for the executive team so that it contains only the data they want to see.
- B. Create a perspective for the executive team.
- C. Create security roles to restrict access to the executive team.
- D. Create a view for the executive team.

Answer: B

Question: 11

You need to configure the permissions for the sales team members in the Sales Reports library.

Which permissions should you use? (Each answer presents part of the solution. Choose all that apply.)

- A. Delete Items
- B. Add Items
- C. View Items
- D. Create Alerts
- E. Manage Alerts
- F. Edit Items

Answer: C, F

Question: 12

You need to configure the SSRS data source.

What should you do?

- A. Use Windows credentials.
- B. Prompt the user for credentials.
- C. In the data source configuration window, select the Credentials are not required option.
- D. Store the credentials.

Answer: A

Question: 13

You need to deploy the StandardReports project.

What should you do? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Use the Analysis Services Deployment utility to create an XMLA deployment script.
- B. Deploy the project from SQL Server Data Tools (SSDT).
- C. Use the Analysis Services Deployment wizard to create an XMLA deployment script.
- D. Use the Analysis Services Deployment wizard to create an MDX deployment script.

Answer: B, C

Question: 14

You need to deploy the StandardReports project at the end of the current business day.

What should you do? (Each correct answer presents a complete solution. Choose all that apply.)

- A. Use the Analysis Services Deployment utility to create an XMLA deployment script and run it at the end of the day.
- B. Use the Analysis Services Deployment wizard to create an MDX deployment script and run it at the end of the day.
- C. Use the Analysis Services Deployment wizard to create an XMLA deployment script and run it at the end of the day.
- D. Deploy the project from SQL Server Data Tools (SSDT) at the end of the day.

Answer: C, D

Case Study: 4,

Mix Questions

Question: 1

You are developing a SQL Server Analysis Services (SSAS) tabular project. The model includes a table named Sales. The Sales table includes a single date column.

The Sales table must meet the following requirements:

Queries must be able to return all rows.

Must be able to support four different processing schedules for different date ranges.

Date ranges must not include any overlapping data.

You need to implement a solution that meets the requirements.

What should you do?

- A. Create four partitions for the Sales table. Create four roles. Use the same row filter queries for each role and partition.
- B. Convert the Sales table into four smaller tables by using row filter queries. Use one perspective for all four tables.
- C. Create four partitions for the Sales table. Use row filter queries for each partition.
- D. Convert the Sales table into four smaller tables by using row filter queries. Use one perspective for each of the four tables.

Answer: C

Question: 2

You are administrating a SQL Server Analysis Services (SSAS) tabular database.

You need to create a new role that allows its members to query data and to refresh data in the model.

Which permission should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Browse and Manage
- B. Administrator
- C. Read and Process
- D. Explore and Manage

Answer: C

Explanation:

* Giving a database role permission to process an Analysis Services database means that the role has permission to perform all processing options on the database. This includes the processing of all cubes, dimensions, mining structures, and mining models in the database. However, the role does not have permission to read database metadata or access any data in the database itself.

Question: 3

You are developing a SQL Server Analysis Services (SSAS) tabular project.

A column named City must be added to the table named Customer. The column will be used in the definition of a hierarchy. The City column exists in the Geography table that is related to the Customer table.

You need to add the City column to the Customer table.

How should you write the calculation?

- A. City:= LOOKUP(Geography[City],Geography[GeographyKey],[GeographyKey])
- B. City:= LOOKUPVALUE(Geography[City],Geography[GeographyKey],[GeographyKey])
- C .=RELATED(Geography[City])
- D. =RELATED(Geography.City)
- E. =VALUES(Geography[City])
- F. City:=VALUES(Geography[City])

Answer: C

Explanation:

* RELATED Function

Returns a related value from another table.

Question: 4

You are managing a SQL Server Analysis Services (SSAS) tabular database.

The database must meet the following requirements:

The processing must load data into partitions or tables.

The processing must not rebuild hierarchies or relationships.

The processing must not recalculate calculated columns.

You need to implement a processing strategy for the database to meet the requirements.

Which processing mode should you use?

- A. Process Clear
- B. Process Data
- C. Process Add
- D. Process Full
- E. Process Default

Answer: C

Question: 5

You are developing a SQL Server Analysis Services (SSAS) tabular project that will be used by the finance, sales, and marketing teams.

The sales team reports that the model is too complex and difficult to use. The sales team does not need any information other than sales-related resources in the tabular model. The finance and marketing teams need to see all the resources in the tabular model.

You need to implement a solution that meets the needs of the sales team while minimizing development and administrative effort.

What should you do?

- A. Create a separate partition for each team.
- B. Create a separate data source for each team.
- C. Create a perspective for the sales team.
- D. Enable client side security to filter non-sales data.

Answer: C

Question: 6

You are developing a SQL Server Analysis Services (SSAS) tabular project.

You need to grant the minimum permissions necessary to enable users to query data in a data model.

Which role permission should you use?

- A. Explorer
- B. Process
- C. Browser
- D. Administrator
- E. Select

F. Read

Answer: F

Question: 7

You are developing a SQL Server Analysis Services (SSAS) tabular project. In the data warehouse, a table named Sales Persons and Territories defines a relationship between a salesperson's name, logon ID, and assigned sales territory. You need to ensure that each salesperson has access to data from only the sales territory assigned to that salesperson. You need to use the least amount of development effort to achieve this goal. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a new role named Sales Persons with Read permission. Add each salesperson's logon as a member to the role.
- B. Add the Sales Persons and Territories table to the model, define the relationships, and then implement dynamic security by using row filters. Grant each salesperson access to the model.
- C. Create a new Active Directory Domain Services (AD DS) security group and add each salesperson as a member. Then create a new role named Sales Persons with Read permission. Add the group as a member to the new role.
- D. Create a separate tabular model for each sales territory and assign each tabular model a corresponding sales territory name. Grant each salesperson access to the corresponding tabular model of the assigned sales territory.

Answer: B

Question: 8

You are troubleshooting query performance for a SQL Server Analysis Services (SSAS) cube. A user reports that a Multidimensional Expressions (MDX) query is very slow. You need to identify the MDX query statement in a trace by using SQL Server Profiler. Which event class should you use?

- A. Get Data From Aggregation
- B. Query Subcube
- C. Query Begin
- D. Progress Report Begin
- E. Calculate Non Empty Begin
- F. Execute MDX Script Begin

Answer: C

Question: 9**DRAG DROP**

You install a SQL Server Analysis Services (SSAS) instance in tabular mode on a server.

While processing a very large tabular model, you receive an out-of-memory error. You identify that the amount of physical memory in the server is insufficient. Additional physical memory cannot be installed in the server.

You need to configure the server to allow paging to disk by using the operating system page file (pagefile.sys).

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Change the value of the **Memory**

VertiPaqPagingPolicy configuration option to **1**.

Change the value of the **OLAP Process**

AllowDiskPaging configuration option to **1**.

Change the value of the **Memory**

VertiPaqPagingPolicy configuration option to **2**.

Restart the Analysis Services instance.

In Object Explorer, right-click the Analysis Services instance and then click **Properties**.

Change the value of the **Memory**

VertiPaqMemoryLimit configuration option to **0**.

Select the **Show Advanced (All)**

Properties checkbox.

Answer:

Box 1:

In Object Explorer, right-click the Analysis Services instance and then click **Properties**.

Box 2:

Select the **Show Advanced (All)**

Properties checkbox.

Box 3:

Change the value of the **Memory**

VertiPaqPagingPolicy configuration option to **1**.

Box 4:

Restart the Analysis Services instance.

Explanation:

Note:

* View or set configuration properties in Management Studio

In SQL Server Management Studio, connect to an Analysis Services instance.

In Object Explorer, right-click the Analysis Services instance, and then click Properties. The General page appears, displaying the more commonly used properties.

To view additional properties, click the Show Advanced (All) Properties checkbox at the bottom of the page.

Modifying server properties is supported only for tabular mode and multidimensional mode servers. If you installed PowerPivot for SharePoint, always use the default values unless you are directed otherwise by a Microsoft product support engineer.

* **VertiPaqPagingPolicy**

Specifies the paging behavior in the event the server runs low on memory. Valid values are as follows:

Zero (0) is the default. No paging is allowed. If memory is insufficient, processing fails with an out-of-memory error.

1 enables paging to disk using the operating system page file (pagefile.sys).

When VertiPaqPagingPolicy is set to 1, processing is less likely to fail due to memory constraints because the server will try to page to disk using the method that you specified. Setting the VertiPaqPagingPolicy property does not guarantee that memory errors will never happen. Out of memory errors can still occur under the following conditions: There is not enough memory for all dictionaries. During processing, Analysis Services locks the dictionaries for each column in memory, and all of these together cannot be more than the value specified for VertiPaqMemoryLimit. There is insufficient virtual address space to accommodate the process.

To resolve persistent out of memory errors, you can either try to redesign the model to reduce the amount of data that needs processing, or you can add more physical memory to the computer.

Applies to tabular server mode only

* Incorrect: VertiPaqMemoryLimit

If paging to disk is allowed, this property specifies the level of memory consumption (as a percentage of total memory) at which paging starts. The default is 60. If memory consumption is less than 60 percent, the server will not page to disk.

This property depends on the VertiPaqPagingPolicyProperty, which must be set to 1 in order for paging to occur.

Applies to tabular server mode only.

Reference: Memory Properties; Configure Server Properties in Analysis Services

Question: 10

You maintain SQL Server Analysis Services (SSAS) instances.

You need to configure an installation of PowerPivot for Microsoft SharePoint in a SharePoint farm.

Which tool should you use? (Each correct answer presents a complete solution. Choose all that apply.)

- A. SQL Server Configuration Manager
- B. PowerPivot Configuration Tool
- C. SharePoint Products Configuration Wizard
- D. SharePoint Central Administration
- E. PowerShell

Answer: A, B, D

Question: 11

You are developing a tabular Business Intelligence Semantic Model (BISM) database based on a SQL Server database. In the data source, the FactInternetSales table is partitioned by month. Data from the current month has been updated and new data has been inserted in the FactInternetSales table, in the DimProduct table, and in the DimCustomer table.

In the model, the FactInternetSales table is also partitioned by month.

You need to ensure that the model has the most recent data while minimizing the processing time.

What should you do?

- A. Process the latest FactInternetSales model table partition, the DimProduct table, and the DimCustomer table with the Process Clear processing option. Then process the database with the Process Data processing option.
- B. Process the latest FactInternetSales model table partition, the DimProduct table, and the DimCustomer table with the Process Clear processing option. Then process the database with the Process Full processing option.
- C. Process the latest FactInternetSales model table partition, the DimProduct table, and the DimCustomer table with the Process Defrag processing option. Then process the database with the Process Recalc processing option.
- D. Process the latest FactInternetSales model table partition, the DimProduct table, and the DimCustomer table with the Process Data processing option. Then process the database with the Process Defrag processing option.
- E. Process the latest FactInternetSales model table partition, the DimProduct table, and the DimCustomer table with

the Process Data processing option. Then process the database with the Process Recalc processing option.

Answer: D

Question: 12

HOTSPOT

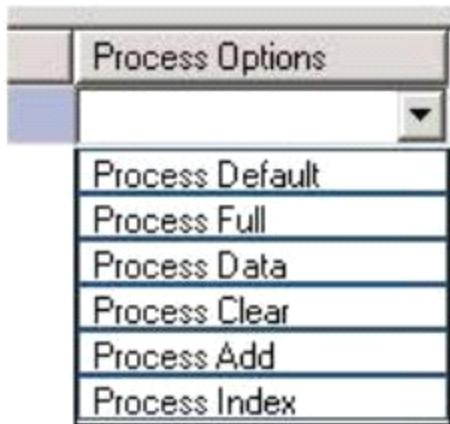
You maintain a multidimensional Business Intelligence Semantic Model (BISM) that was developed with default settings.

The model has one cube and the cube has one measure group. The measure group is based on a very large fact table and is partitioned by month. The fact table is incrementally loaded each day with approximately 800,000 new rows.

You need to ensure that all rows are available in the cube while minimizing the processing time.

Which processing option should you use? (To answer, select the appropriate option in the answer area.)

Work Area



Answer:

Explanation:

Process Add

Question: 13

You are modifying a SQL Server Analysis Services (SSAS) multidimensional database.

You have identified a dimension that is no longer used by any cubes.

You need to delete the dimension.

What should you do?

- A. Write a Multidimensional Expressions (MDX) command to drop the dimension from the database.
- B. Write a Data Mining Extensions (DMX) command to drop the dimension from the database.
- C. Script the deletion of the dimension as an XMLA command for execution against the production model.
- D. Write a T-SQL command to drop the dimension from the database.

Answer: C

Question: 14

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) multidimensional project that is configured to source data from a Microsoft Azure SQL Database database. The cube is processed each night at midnight.

The largest partition in the cube takes 12 hours to process, and users are unable to access the cube until noon. The partition must be available for querying as soon as possible after processing commences.

You need to ensure that the partition is available for querying as soon as possible, without using source data to satisfy the query.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Open the cube for editing and then select the **Partitions** tab.

On the Partitions tab, select the partition to edit.

Enable proactive caching and then select the **Bring Online Immediately** option.

Click **Storage Settings** and then click **Options** to open the Storage Options dialog box.

On the Properties window, change the **ProcessingMode** property to **LazyAggregations**.

On the Properties window, change the **ProcessingMode** property to **(default)**.

Answer:

Box 1:

Open the cube for editing and then select the **Partitions** tab.

Box 2:

On the Partitions tab, select the partition to edit.

Box 3:

On the Properties window, change the **ProcessingMode** property to **LazyAggregations**.

Explanation:

Note:

* Processing mode has two possible options.

Regular. This is the default setting. When set to regular, partitions will be available to users after data has been loaded and aggregations are created completely.

Lazy Aggregations. When set to lazy aggregations, partitions will be available to user queries immediately after data has been loaded. Aggregations will be created as a separate background process while users start to query the partition.

* Lazy processing performs the task of building indexes and aggregations for dimensions and measure group partitions at a lower priority to reduce foreground processing time and to allow users to query the cube sooner. For lazy processing to occur, you must switch the **ProcessingMode = LazyAggregations** of your measure group partitions; by default this value is Regular (lazy processing is turned off). When processing a dimension with flexible aggregations such as parent-child or virtual dimension by using the processing enumeration of **ProcessUpdate** (such as to take into account of member name or hierarchy changes), lazy processing is initiated to ensure that the aggregations are rebuilt on the associated measure group partitions.

* Configure Lazy Processing for the cube, measure group, or partition. If you configure Lazy Processing, the dropped aggregations are recalculated as a background task. While the flexible aggregations are being recalculated, users can continue to query the cube (without the benefit of the flexible aggregations). While the flexible aggregations are being recalculated, queries that would benefit from the flexible aggregations run slower because Analysis Services resolves these queries by scanning the fact data and then summarizing the data at query time. As the flexible

aggregations are recalculated, they become available incrementally on a partition-by-partition basis. For a given cube, Lazy Processing is not enabled by default. You can configure it for a cube, measure group, or partition by changing the ProcessingMode property from Regular to LazyAggregations. To manage Lazy Processing, there are a series of server properties such as the LazyProcessing \ MaxObjectsInParallel setting, which controls the number of objects that can be lazy processed at a given time. By default it is set to 2. By increasing this number, you increase the number of objects processed in parallel; however, this also impacts query performance and should therefore be handled with care.

* Incorrect: With Bring Online Immediately enabled, during cache refresh all queries are directed to the relational source database to retrieve the latest data for end users. While this provides users with refreshed data, it can also result in reduced query performance given that Analysis Services needs to redirect queries to the relational source database.

Question: 15

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) multidimensional project.

You identify that a dimension, which has a large number of attributes, is issuing a separate (SELECT DISTINCT) query for each attribute even though they all come from a single database table. The table does not have a large number of rows.

You need to configure the processing of the dimension to issue only a single SQL query to the underlying database, and processing must continue if any errors are encountered.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Actions	Answer area
In the Properties window, set the ProcessingGroup property to ByTable and then change the KeyDuplicate property to ReportAndStop .	
In the Properties window, set the ProcessingMode property to ByTable and then change the KeyDuplicate property to IgnoreError .	
Open the dimension designer.	
In the Properties window, set the ProcessingMode property to ByTable and then change the KeyDuplicate property to ReportAndStop .	
Select the dimension.	
In the Properties window, set the ProcessingGroup property to ByTable and then change the KeyDuplicate property to IgnoreError .	
Open the cube designer.	

Answer:

Box 1:

Open the dimension designer.

Box 2:

Select the dimension.

Box 3:

In the Properties window, set the **ProcessingGroup** property to **ByTable** and then change the **KeyDuplicate** property to **IgnoreError**.

Question: 16

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) tabular project based on a Microsoft Azure SQL Database database. The ProcessingOption property for the project is set to Do Not Process.

Several calculated columns have been added to a table. The project has been deployed to the production server.

You need to ensure that newly added data is processed on the production server.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Right-click the table and then select **Process Table**.

Open the project in SQL Server Data Tools (SSDT).

On the **Model** menu, select **Process** and then select **Process Table**.

In the model designer, select the table.

In Object Explorer, connect to the SSAS instance, expand the database, and then expand the Tables folder.

Open SQL Server Management Studio (SSMS).

Answer:

Box 1:

Open SQL Server Management Studio (SSMS).

Box 2:

In Object Explorer, connect to the SSAS instance, expand the database, and then expand the Tables folder.

Box 3:

Right-click the table and then select **Process Table**.

Explanation:

Note:

* To process a table

In SQL Server Management Studio, in the tabular model database which contains the table you want to process, expand the Tables node, then right-click on the table you want to process, and then click Process Table.

In the Process Table dialog box, in the Model list box, select one of the following process modes:

Etc-

Reference: Process Database, Table, or Partition

Question: 17**HOTSPOT**

A SQL Server Analysis Services (SSAS) cube has roles to define dimension data security. A role named USA allows users to browse data pertaining to the United States. A role named Canada allows users to browse data pertaining to Canada.

A user can browse sales data pertaining to the United States but cannot browse sales data pertaining to Canada. You validate that the user belongs to the USA and Canada roles.

You need to reproduce the issue in SQL Server Management Studio (SSMS).

Which option should you select? To answer, select the appropriate action in the answer area.

Country	Reseller Order Quan...
Australia	3009
Canada	18801
France	7715
Germany	4480
United Kingdom	7060
United States	59107

Answer:

Country	Reseller Order Quan...
Australia	3009
Canada	18801
France	7715
Germany	4480
United Kingdom	7060
United States	59107

Question: 18

You manage an environment that has SharePoint Server 2010 and a SQL Server Reporting Services (SSRS) instance in SharePoint integrated mode. Several report subscriptions are configured to deliver reports through email by using a shared schedule.

The email server will be going offline.

You need to temporarily suspend the shared schedule until the email server is brought back online.

What should you do?

- A. In Report Manager, pause the shared schedule.
- B. In SharePoint Central Administration, pause the shared schedule.
- C. In Report Manager, delete the shared schedule.
- D. In SharePoint Central Administration, delete the shared schedule.

Answer: B

Question: 19

DRAG DROP

You manage a SQL Server Reporting Services (SSRS) instance in native mode. You are building a shared dataset for your weekly performance reports. The shared dataset uses a data source that is configured to use credentials that are stored in the Report Server.

You have a predefined shared schedule to perform cleanup and maintenance tasks for SSRS.

You need to enable caching on the shared dataset. You also need to use an existing shared schedule to discard the cache.

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

- Open SharePoint Central Administration and then click the shared dataset.
- Select the Caching page and then click the **Cache shared dataset** checkbox.
- Open Report Manager and then click the shared dataset.
- Select the **Expire the cache on the following schedule** option and then select the **Shared Schedule** option.
- From the combo box, select the shared schedule and then click **Apply**.

Answer:

Box 1: Open Report Manager and then click the shared dataset.

We should use Report Manager (not SharePoint Central Administration).

Box 2: Select the Caching page and then click the Cache shared dataset checkbox.

We set up caching.

Box 3: Select the Expire the cache on the following schedule option and then select the Shared Schedule option.

We configure caching further.

Box 4: From the combo box, select the shared schedule and then click Apply.

Finally we define scheduling.

Explanation:

Note on caching:

To open the Caching properties page for a shared dataset

Open Report Manager, and locate the report for which you want to configure shared dataset properties.

Point to the shared dataset, and click the drop-down arrow.

In the drop-down list, click Manage. The General properties page for the report opens.

Click the Caching tab.

Options include:

Cache shared dataset (Box 2 above)

Places a temporary copy of the data in a cache when a user first opens a report that uses this shared dataset. Subsequent users who run the report within the caching period receive the cached copy of the data. Caching usually improves performance because the data is returned from the cache instead of running the dataset query again.

Expire the cache on the following schedule (box 3 above)

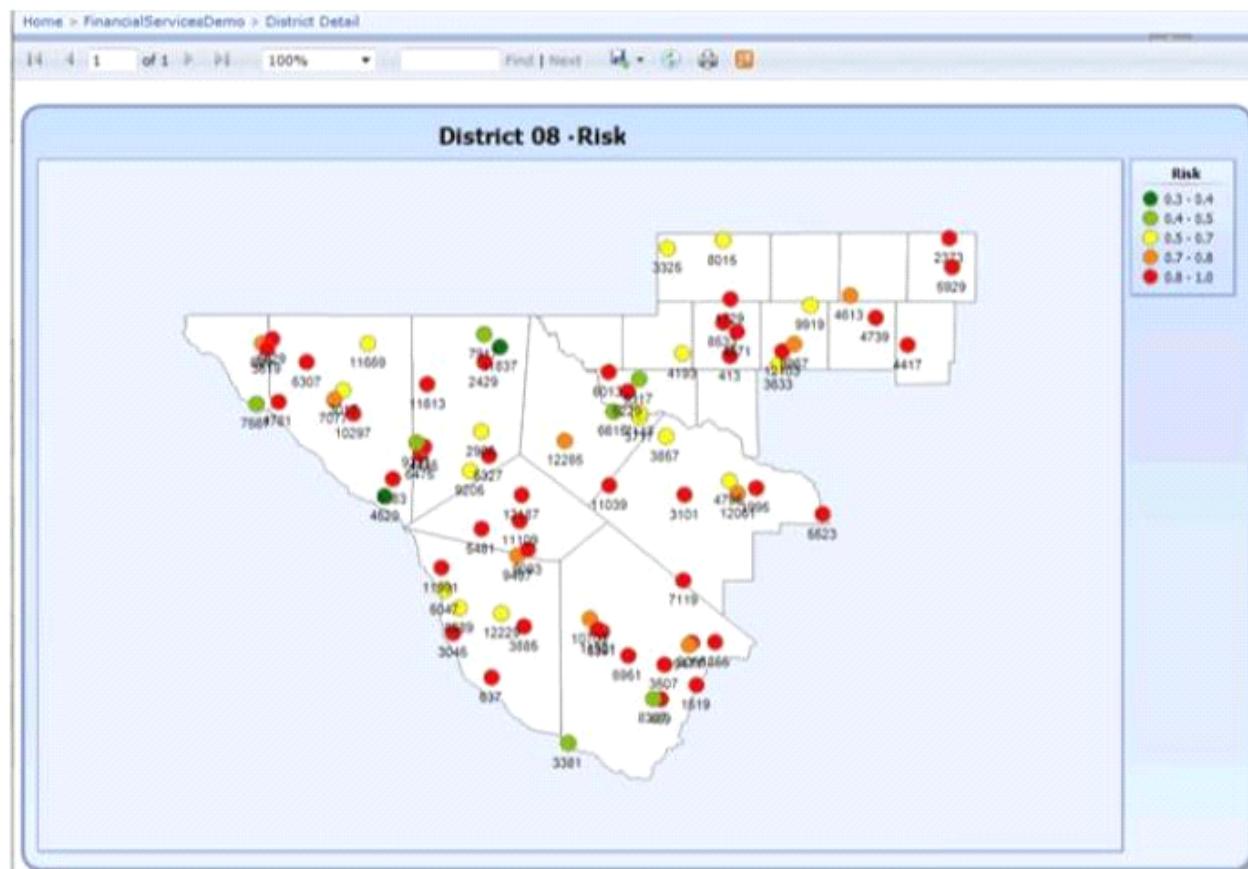
Schedule the time when the cached data is no longer valid and is removed from the cache. The schedule can be a shared schedule or one that is specific for only the current shared dataset.

Reference: SQL Server 2012, Caching Page, Shared Datasets (Report Manager)

Question: 20

You are designing a SQL Server Reporting Services (SSRS) report for a bank. The bank has Automated Teller Machines (ATMs) in several regions. ATM operational data is stored in a Microsoft Azure SQL Database database.

The report must use a map to display the location and status of the ATMs as shown in the following exhibit. (Click the Exhibit button.)



You need to ensure that the report displays only a user selected map region.

Which source of spatial data should you use for the map?

A. SQL Server spatial query

B. Map gallery

- C. ESRI shape file
- D. Bing Maps layer

Answer: A

Question: 21

You are developing a SQL Server Reporting Services (SSRS) report. The report includes a dataset with fields named Year Month Number, and InvCount. The report includes a table that displays the inventory count per year, as shown in the following table.

Year	Inv Count
1995	46,043
1996	45,471
1997	45,765
1998	45,484
1999	47,193
2000	48,456
2001	48,412
2002	47,902
2003	48,049
2004	48,442
2005	48,519
2006	48,837
2007	49,074
2008	48,981
2009	49,251
2010	49,407
2011	49,547

You need to modify the table to include a graphical item displaying the inventory count trend to the right of the Inv Count column.

What should you do?

- A. Add an Indicator item to a new column on the right of the Inv Count column. Select the Directional Indicator type and then assign the MonthNumber field to the Start property.
- B. Add a Sparkline item to a new column on the right of the Inv Count column. Then select the InvCount field for Values and the MonthNumber field for Series Groups.
- C. Add a Sparkline item to a new column on the right of the Inv Count column. Then select the InvCount field for Values and the MonthNumber field for Category Groups.
- D. Add a text box to a new column on the right of the Inv Count column. Then use a Go to report action to link to a separate report showing the monthly trend.
- E. Add an Indicator item to a new column on the right of the Inv Count column. Select the Directional Indicator type and then select the MonthNumber field for Value.

Answer: C

Question: 22

You are developing a SQL Server Reporting Services (SSRS) report for a company that has a subscription to a dataset

from the Microsoft Azure Data Market.

The dataset permits flexible queries.

You need to choose the connection type for the data source.

Which connection type should you choose?

- A. XML
- B. Microsoft Azure SQL Database
- C. Microsoft SharePoint List
- D. ODBC

Answer: A

Question: 23

You are designing a SQL Server Reporting Services (SSRS) report.

The report defines a single SQL Server data source and dataset.

You need to include additional data sourced from a Microsoft Azure SQL Database in the report.

What should you do?

- A. Create a SQL Server data source and then add a dataset that uses the new data source.
- B. Create a Microsoft Azure SQL Database data source and then add a dataset that uses the new data source.
- C. Generate an Atom-compliant data feed for the report.
- D. Create a Microsoft Azure SQL Database dataset that uses the existing data source.

Answer: B

Question: 24

You are developing a new SQL Server Reporting Services (SSRS) report in SQL Server Data Tools (SSDT). This report has a table named Table1 and a textbox named Textbox1.

Table1 is initially visible but the user must be able to choose when to hide it.

You need to develop the report to meet the requirement.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. For the properties of Table1, configure the Display can be toggled by this report item option to use Textbox1.
- B. Configure Textbox1 to drill through to rerun the report to toggle the display of Table1.
- C. For the properties of Table1, configure the Display can be toggled by this report item option to use Table1
- D. Add a parameter to the report so users can choose the display state of Table1.

Answer: A

Question: 25

You are developing a new SQL Server Reporting Services (SSRS) report in SQL Server Data Tools (SSDT).

The report must define a report parameter to prompt the user for the business unit. Each business unit has a unique font scheme combination of font and size properties.

You need to ensure that all of the text boxes in the table headers use the correct business unit font properties.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Add one report variable for FontFamily Assign it with an expression to return the appropriate colors. For each header text box, set the Color and FontSize properties by using the variables.
- B. For each header text box, assign expressions to the FontFamily and FontSize properties.
- C. Add two report variables named FontFamily and FontSize. Assign them with expressions to return the appropriate colors. For each header text box, use expressions to set the FontFamily and FontSize properties by using the variables.
- D. Add two Microsoft Visual C# functions to the code block of the report to implement FontFamily and FontSize functions. For each header text box, use expressions to set the FontFamily and FontSize properties by using the functions.

Answer: D

Question: 26

You install SQL Server Reporting Services (SSRS).

You need to restore a copy of the symmetric key.

Which command should you run?

- A.rskeymgmt -d
- B.rskeymgmt -e -f %temp%\rs.key -p Password1
- C.rskeymgmt -i
- D.rskeymgmt -a -f %temp%\rs.key -p Password1

Answer: D

Question: 27

You are managing a SQL Server Reporting Services (SSRS) instance in native mode. A role named Folder Access Controller is present on the server.

The Folder Access Controller role consists of only the Set security for individual items task.

When role members open Report Manager, they cannot view folders.

You need to modify the Folder Access Controller role so that the role members can view folders.

Which task should you add to the Folder Access Controller role?

- A. Manage models
- B. Manage reports
- C. View reports
- D. Manage folders

Answer: D

Question: 28

You have a SQL Server Reporting Services (SSRS) instance. The instance has a report that displays 3 million records.

Users report that they experience performance issues when they use the report.

You need to ensure that the report renders as quickly as possible.

Which two actions should you perform? Each correct answer presents part of the solution.

- A. Enable versioning.
- B. Enable caching.
- C. Enable report history.

- D. Create snapshots.
- E. Implement pagination.

Answer: A, B

Question: 29

You are working with a SQL Server Reporting Services (SSRS) instance in native mode. An item role named Reports Writer is present on the server.

The Reports Writer role cannot view and modify report caching parameters.

You need to ensure that the Reports Writer role can view and modify report caching parameters.

What should you do?

- A. Add the Manage all subscriptions task to the Reports Writer role.
- B. Add the Manage report history task to the Reports Writer role.
- C. Add the View data sources task to the Reports Writer role.
- D. Add the Manage individual subscriptions task to the Reports Writer role.

Answer: B

Question: 30

You are developing a SQL Server Analysis Services (SSAS) multidimensional project.

A fact table is related to a dimension table named DimScenario by a column named ScenarioKey.

The dimension table contains three rows for the following scenarios:

Actual

Budget Q1

Budget Q3

You need to create a dimension to allow users to view and compare data by scenario.

What should you do?

- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- F. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- G. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- H. Create several dimensions. Add each dimension to the cube.
- I. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- J. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- K. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.

- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: K

Question: 31

You are developing a SQL Server Analysis Services (SSAS) cube.

The data warehouse has a table named FactStock that is used to track movements of stock. A column named MovementQuantity contains quantities of stock. A positive quantity is used for input and negative quantity is used for output. A column named Movement Date is related to the time dimension. The quantity in stock, at a given point in time, can be evaluated as the sum of all MovementQuantity values at that point in time.

You need to create a measure that calculates the quantity in stock value.

What should you do?

- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- F. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- G. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- H. Create several dimensions. Add each dimension to the cube.
- I. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- J. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- K. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: G

Question: 32

You are designing a SQL Server Analysis Services (SSAS) cube.

You need to create a measure to count unique customers.

What should you do?

- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- F. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- G. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- H. Create several dimensions. Add each dimension to the cube.
- I. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- J. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- K. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: E

Question: 33

You are creating a SQL Server Analysis Services (SSAS) cube.

You need to create a time dimension. It must be linked to a measure group named Sales at the day granularity level. It must also be linked to a measure group named Salary at the month granularity level.

What should you do?

- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- F. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- G. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- H. Create several dimensions. Add each dimension to the cube.
- I. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- J. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- K. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the

- DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: J

Question: 34

You are developing a SQL Server Analysis Services (SSAS) cube for the sales department at your company.

The sales department requires the following set of metrics:

Unique count of customers

Unique count of products sold

Sum of sales

You need to ensure that the cube meets the requirements while optimizing query response time.

What should you do? (Each answer presents a complete solution. Choose all that apply.)

- A. Place the measures in a single measure group.
- B. Place the distinct count measures in separate measure groups.
- C. Use the additive measure group functions.
- D. Use the semiadditive measure group functions.
- E. Use the Count and Sum measure aggregation functions.
- F. Use the Distinct Count and Sum measure aggregation functions.

Answer: B, F

Question: 35

DRAG DROP

You are developing reports based on the SQL Server Analysis Services (SSAS) cube named ProcessedOrders.

A Multidimensional Expressions (MDX) query must include a query-scoped calculated member, which computes average sales per order item. The query must also return the set of three states in a query-scoped named set named East Coast Customers.

You need to define the calculations in an MDX query to meet the requirements.

Which four MDX segments should you insert in sequence before a SELECT statement? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

```
MEMBER [Average Sales per Processed Order  
Item] AS  
  
MEASURE [Average Sales per Processed Order  
Item] AS  
  
WITH SET [East Coast Customers] AS  
  
{ [Measures].[Sales]/[Measures].[ProcessedOrder  
Quantity] }  
  
{ [Measures].[Sales]/[Measures].[ProcessedOrder  
Quantity] }  
  
;  
  
{ [Customer].[Geography].[State].&[QLD]&[AU],  
[Customer].[Geography].[State].&[NSW]&[AU],  
[Customer].[Geography].[State].&[VIC]&[AU] }
```

Answer:

Box 1:

WITH SET [East Coast Customers] AS

Box 2:

```
{[Customer].[Geography].[State]&[QLD]&[AU],  
[Customer].[Geography].[State]&[NSW]&[AU],  
[Customer].[Geography].[State]&[VIC]&[AU]}
```

Box 3:

```
MEMBER [Average Sales per Processed Order  
Item] AS
```

Box 4:

```
([Measures].[Sales]/[Measures].[ProcessedOrder  
Quantity])
```

Explanation:

Note:

* Example #1:

WITH

```
MEMBER [Measures].[Special Discount] AS
```

```
[Measures].[Discount Amount] * 1.5
```

SELECT

```
[Measures].[Special Discount] ON COLUMNS,
```

```
NON EMPTY [Product].[Product].MEMBERS ON Rows
```

FROM [Adventure Works]

WHERE [Product].[Category].[Bikes]

* Example 2:

WITH SET [ChardonnayChablis] AS

```
{[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Good].[Good Chardonnay],  
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Pearl].[Pearl Chardonnay],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Portsmouth].[Portsmouth  
Chardonnay],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Top Measure].[Top Measure  
Chardonnay],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Walrus].[Walrus Chardonnay],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Good].[Good Chablis Wine],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Pearl].[Pearl Chablis Wine],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Portsmouth].[Portsmouth Chablis  
Wine],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Top Measure].[Top Measure Chablis  
Wine],
```

```
[Product].[All Products].[Drink].[Alcoholic Beverages].[Beer and Wine].[Wine].[Walrus].[Walrus Chablis Wine]}
```

SELECT

```
[ChardonnayChablis] ON COLUMNS,
```

```
{Measures.[Unit Sales]} ON ROWS
```

FROM Sales

Reference:

<http://technet.microsoft.com/en-us/library/ms146017.aspx>

<http://technet.microsoft.com/en-us/library/ms145487.aspx>

Question: 36

You work in the Business Intelligence (BI) department of a multinational company.

The company has requested a new corporate BI solution that meets the following requirements:

- The solution must use SQL Server Analysis Services (SSAS).
- The model must incrementally add 10 million fact rows per month.
- The model must be translated to English, French, or Spanish based on users' locale.
- The model must be able to contain the most recent 36 months of data.

You need to select the appropriate model type and partitioning strategy to meet the requirements.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a tabular model with one partition for all of the data.
- B. Create a multidimensional model with one partition for each month.
- C. Create a multidimensional model with one partition for all of the data.
- D. Create a tabular model with one partition for each month.

Answer: B

Question: 37

You are developing a BI Semantic Model (BISM) that will be used to analyze complex budgeting and forecast data sourced from a financial database. The model will be deployed to a server with 28 GB of RAM.

The source data, located in a SQL Server data warehouse, is currently using 15 terabytes of disk space and is doubling in size every month. The model will be queried by staff in the accounting department by using Microsoft Excel 2010.

You need to ensure the highest query performance and scalability for the accounting department queries.

Which project type should you choose?

- A Tabular project that uses the In-Memory query mode
- B PowerPivot workbook deployed to SharePoint
- C Tabular project that uses the DirectQuery query mode
- D Multidimensional project

Answer: D

Question: 38

HOTSPOT

You have a database named DB1. DB1 contains three tables named FactSales, DimCustomer, and DimEmployee. A sample from DimEmployee is shown in the following table.

EmployeeID	Name	Title	Managers
1	Kim Abercrombie	CEO	<i>Not applicable</i>
20	Brad Sutton	Sales Manager	1
31	Don Funk	Sales Supervisor	20
123	Diane Margheim	Sales Associate	31

A sample from DimCustomer is shown in the following table.

CustomerID	Name	City	State	Country
1	Contoso, Ltd.	Miami	FL	United States
2	Adventure Works	Montreal	QC	Canada
3	Woodgrove Bank	London	<i>Not applicable</i>	United Kingdom

A sample from FactSales is shown in the following table.

ID	Customer	Employee	Date	Amount
21	1	20	1/1/2012	2345.67
27	2	123	1/1/2012	23254.45
34	3	31	1/2/2012	43543.56

You create a cube named Sales that uses the three tables. FactSales is used as a fact table, DimCustomer is used to create a dimension named Customer, and DimEmployee is used to create a dimension named Employee.

Users report the following issues:

When browsing the Customer dimension, the users see a blank member in the state hierarchy under United Kingdom.

When browsing the Employee dimension, the sales data for Sales Managers and Sales Supervisors is NOT displayed.

You need to resolve the issues by configuring the properties of the dimension attributes.

What should you do? To answer, select the appropriate value for each property in the answer area.

Answer Area

HideMemberIf property for State:

HideMemberIf property for City:

MembersWithData property for EmployeeID:

Answer Area

HideMemberIf property for State:
 Never
 NoName
 OnlyChildWithNoName
 OnlyChildWithParentName
 ParentName

HideMemberIf property for City:
 Never
 NoName
 OnlyChildWithNoName
 OnlyChildWithParentName
 ParentName

MembersWithData property for EmployeeID:
 NonLeafDataHidden
 NonLeafDataVisible

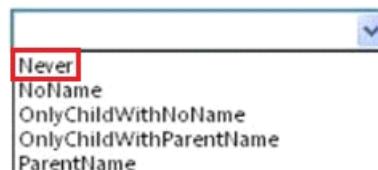
Answer:

Answer Area

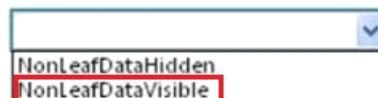
HideMemberIf property for State:



HideMemberIf property for City:



MembersWithData property for EmployeeID:

**Question: 39**

You are developing a BI Semantic Model (BISM) based on a simple and small dataset sourced from SQL Server. The data size and complexity of the data relationships will not change. The model will be used to produce reports in Power View.

You need to use an appropriate project type.

Which project types should you use? (Each answer presents a complete solution. Choose all that apply.)

- A. A tabular project that uses the In-Memory query mode
- B. A tabular project that uses the DirectQuery query mode
- C. A multidimensional project that uses the MOLAP storage mode
- D. A PowerPivot workbook that is deployed to Microsoft SharePoint Server 2010
- E. A multidimensional project that uses the ROLAP storage mode

Answer: A, B, D

Power View is a thin web client that launches right in the browser from a data model in SharePoint Server 2010. The model can be a PowerPivot model workbook or a tabular model running on a SQL Server 2012 Analysis Services (SSAS) server.

Question: 40**HOTSPOT**

You are designing a SQL Server Analysis Services (SSAS) cube that contains two measure groups named Sales History and Current Sales.

The Sales History measure group has the following characteristics:

Data changes on a monthly basis.

The measure group contains a very high data volume.

Queries that use the measure group only reference aggregate data.

The measure group includes SUM, MIN, MAX, and COUNT aggregate functions.

The Current Sales measure group has the following characteristics:

Data changes frequently.

The measure group contains a low data volume.

Queries that use the measure group often reference non-aggregate data.

The measure group includes SUM, MIN, MAX, and COUNT aggregate functions.

You need to select a storage mode for each measure group. The solution must meet the following requirements:

For the Sales History measure group, query performance must be optimized over data latency.

For the Current Sales measure group, data latency must be optimized over query performance.

What should you do? To answer, select the appropriate storage mode for each measure group in the answer area.

Answer Area

Sales History	Current Sales
<input type="text"/>	<input type="text"/>

Answer Area

Sales History	Current Sales
<input type="text"/> <ul style="list-style-type: none"> Real-time ROLAP Real-time HOLAP Low-latency MOLAP Medium-latency MOLAP Automatic MOLAP Scheduled MOLAP MOLAP 	<input type="text"/> <ul style="list-style-type: none"> Real-time ROLAP Real-time HOLAP Low-latency MOLAP Medium-latency MOLAP Automatic MOLAP Scheduled MOLAP MOLAP

Answer:

Answer Area

Sales History	Current Sales
<input type="text"/> <ul style="list-style-type: none"> Real-time ROLAP Real-time HOLAP Low-latency MOLAP Medium-latency MOLAP Automatic MOLAP Scheduled MOLAP MOLAP 	<input type="text"/> <ul style="list-style-type: none"> Real-time ROLAP Real-time HOLAP Low-latency MOLAP Medium-latency MOLAP Automatic MOLAP Scheduled MOLAP MOLAP

Question: 41

You are modifying a SQL Server Analysis Services (SSAS) cube.

Users of the cube report that the precision for the SalesAmount measure is four digits. You need to ensure that the SalesAmount measure stores values to two digits of precision.

What should you do?

- A. Use the FormatString measure property to format SalesAmount as #,##0.00;-#,##0.00
- B. Use the MeasureExpression measure property to change the precision of SalesAmount to two digits.
- C. Use the FormatString measure property to format SalesAmount as Currency.
- D. Add a named query in the data source view that casts the data source column to two digits of precision. Bind the SalesAmount measure to the new query.
- E. Add a named calculation in the data source view that casts the data source column to two digits of precision. Bind the SalesAmount measure to the new column.

Answer: E

Question: 42

You are developing a multidimensional project that includes a dimension named Organization. The dimension is based on the DimOrganization table in the data warehouse. The following diagram illustrates the table design.

DimOrganization	
PK	<u>OrganizationKey</u>
FK2	ParentOrganizationKey PercentageOfOwnership OrganizationName ParentOrganizationName CurrencyKey
FK1	

The Organization dimension includes a parent-child hierarchy named Organizations. The dimension includes the following dimension attributes:

Organization, which is a key attribute

Organizations, which defines the parent-child hierarchy

Currency Code, which is a regular attribute

PercentageOfOwnership, which is a regular attribute

When users browse the dimension, four hierarchies are visible to them.

You need to ensure that the Organization and PercentageOfOwnership hierarchies are not visible to users.

What should you do?

- A. Set the AttributeHierarchyVisible property to False for the Organization and PercentageOfOwnership attributes.
- B. Set the AttributeHierarchyEnabled property to False for the Organization and PercentageOfOwnership attributes.
- C. Delete the Organization and the PercentageOfOwnership attributes.
- D. Set the AttributeHierarchyDisplayFolder property to Null for the Organization and PercentageOfOwnership attributes.

Answer: A

Question: 43

DRAG DROP

You are developing a SQL Server PowerPivot model to monitor the rate of production of defective products. The PowerPivot model contains many tables. The FactProduction table is configured as shown in the following diagram.



The model requires a key performance indicator (KPI) named Defective Units to measure the rate of production of defective units. The rate of production of defective units is calculated by dividing the sum of the Defects column by the sum of the Units column.

The KPI has the following requirements.

If the rate of production of defective products equals or exceeds 5% then the KPI must report a status of off target (red).

If the rate of production of defective products is less than 5% and equals or exceeds 2% then the KPI must report a status of slightly off target (yellow).

If the rate of production of defective products less than 2% then the KPI must report a status of on target (green).

You need to develop the KPI to meet the requirements.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Actions	Answer area
Configure the KPI to use the Defective Units measure as the target value. Select the Red-Yellow-Green threshold bands. Configure the threshold values to 2% and 5%.	
Create a KPI based on the Units Produced measure.	
Create a KPI based on the Defective Units measure.	
Create a measure named Defective Units to divide the sum of the Defects column by the sum of the Units column.	
Create a measure named Units Produced to sum the Units column. Then create a measure named Defective Units to sum the Defects column.	
Configure the KPI to use the Units Produced measure as the target value. Select the Green-Yellow-Red threshold bands. Configure the threshold values to 2% and 5%.	
Configure the KPI to use the Defective Units measure as the target value. Select the Green-Yellow-Red threshold bands. Configure the threshold values to 2% and 5%.	
Configure the KPI to use the Units Produced measure as the target value. Select the Red-Yellow-Green threshold bands. Configure the threshold values to 2% and 5%.	

Answer:

Box 1:

Create a measure named **Defective Units** to divide the sum of the **Defects** column by the sum of the **Units** column.

Box 2:

Create a KPI based on the **Defective Units** measure.

Box 3:

Configure the KPI to use the **Defective Units** measure as the target value. Select the Green-Yellow-Red threshold bands. Configure the threshold values to 2% and 5%.

Note:

* (Step 1, step 2): The model requires a key performance indicator (KPI) named Defective Units to measure the rate of production of defective units. The rate of production of defective units is calculated by dividing the sum of the Defects column by the sum of the Units column.

* Step 3: The KPI has the following requirements.

- If the rate of production of defective products equals or exceeds 5% then the KPI must report a status of off target (red).
- If the rate of production of defective products is less than 5% and equals or exceeds 2% then the KPI must report a status of slightly off target (yellow).
- If the rate of production of defective products less than 2% then the KPI must report a status of on target (green).

Question: 44

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) cube.

You need to reuse a measure group from a different database.

In SQL Server Data Tools (SSDT), which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

From the **Select a Data Source** step, reference the Analysis Services data source.

From the **Select Objects** step, select the measure group and the dimensions that you need to link.

Launch the Linked Object Wizard.

Launch the Business Intelligence Wizard.

From the **Select Objects** step, select only the measure group that you need to link.

Answer:

Box 1:

Launch the Linked Object Wizard.

Box 2:

From the Select a Data Source step, reference the Analysis Services data source.

Box 3:

From the Select Objects step, select only the measure group that you need to link.

Note:

- * You can use the Linked Object Wizard to either link to or import cubes, dimensions, measure groups, calculations, and Key Performance Indicators (KPIs). You can link to or import these items from another database on the same server or from a database on a remote server
- * On the Select a Data Source page of the Linked Object Wizard, choose the Analysis Services data source or create a new one.
- * On the Select Objects page of the wizard, choose the dimensions you want to link to in the remote database. You cannot link to linked dimensions in the remote database.

* Incorrect:

The Business Intelligence Wizard can guide you through some or all the following steps:

Define time intelligence for cubes.

Define account intelligence for cubes and dimensions.

Define dimension intelligence for cubes and dimensions.

Define unary operators for cubes.

Set custom member formulas for cubes and dimensions.

Specify attribute ordering for dimensions.

Enable dimension writeback for dimensions.

Define semi-additive behavior for cubes.

Define currency conversion for cubes.

Reference: Using Linked Objects in a Cube

Question: 45

HOTSPOT

You are developing a SQL Server Analysis Services (SSAS) cube.

You are writing the following Multidimensional Expressions (MDX) statement for use by a calculated measure. The measure computes the sales amount for the same time period of the previous year. (Line numbers are included for reference only.)

```

01 CREATE MEMBER CURRENTCUBE.Measures.SamePeriodPreviousYearSales AS
02 (Measures.[Sales Amount],
03 (
04     [Date Order].[Calendar].[Calendar Year],
05     1,
06     [Date Order].[Calendar].CurrentMember)),
07 FORMAT_STRING = "#,#.00";

```

You need to complete the MDX statement.

Which MDX function should you use in line 03? To answer, select the appropriate MDX function in the functions list.



Answer:



Question: 46

You are developing a SQL Server Analysis Services (SSAS) multidimensional database. The underlying data source does not have a time dimension table.

You need to implement a time dimension.

What should you do?

- A. Use the SQL Server Data Tools Dimension Wizard and generate a time table on the server.
- B. Create a CSV file with time data and use the DMX IMPORT statement to import data from the CSV file.
- C. Create a time dimension by using the Define dimension intelligence option in the Business Intelligence Wizard.
- D. Create a time dimension by using the Define time intelligence option in the Business Intelligence Wizard.

Answer: A

Question: 47

You are developing a SQL Server Analysis Services (SSAS) cube. The cube contains several dimensions, a local measure group, and a linked measure group. Both measure groups use MOLAP partitions.

You need to write-enable one of the linked measure group partitions to support Microsoft Excel 2010 PivotTable What-If Analysis.

What should you do before the partition can be write-enabled?

- A. Set the Type property of the partition's measure group to Forecast.
- B. Implement the linked measure group as a local measure group.
- C. Implement the local measure group as a linked measure group.
- D. Set the StorageMode property of the linked measure group to Rolap.

Answer: B

Question: 48

DRAG DROP

You have a database named DB1. DB1 contains four tables named FactSales, DimTime, DimCustomer, and DimEmployee. A sample from DimEmployee is shown in the following table.

EmployeeID	Name	Title	Managers
1	Kim Abercrombie	CEO	<i>Not applicable</i>
20	Brad Sutton	Sales Manager	1
31	Don Funk	Sales Supervisor	20
123	Diane Margheim	Sales Associate	31

A sample from DimCustomer is shown in the following table.

CustomerID	Name	City	State	Country
1	Contoso, Ltd.	Miami	FL	United States
2	Adventure Works	Montreal	QC	Canada
3	Woodgrove Bank	London	<i>Not applicable</i>	United Kingdom

A sample from FactSales is shown in the following table.

ID	Customer	Employee	ShipDate	DeliveryDate	Amount
21	1	20	1	8	2345.67
27	2	123	1	8	23254.45
34	3	31	2	9	43543.56

You need to identify which type of dimension must be created for each table.

Which type of dimension should you create for each table? To answer, drag the appropriate dimensions to the correct tables. Each dimension may be used once, more than once, or not at all.

You may need to drag the split bar between panes or scroll to view content.

Dimensions		Answer Area						
<input type="checkbox"/> Degenerate dimension <input type="checkbox"/> Ragged dimension <input type="checkbox"/> Role-playing dimension <input type="checkbox"/> Parent-child dimension		<table border="1"> <tr> <td>DimCustomer</td> <td>Dimension</td> </tr> <tr> <td>DimEmployee</td> <td>Dimension</td> </tr> <tr> <td>DimTime</td> <td>Dimension</td> </tr> </table>	DimCustomer	Dimension	DimEmployee	Dimension	DimTime	Dimension
DimCustomer	Dimension							
DimEmployee	Dimension							
DimTime	Dimension							

Answer:

DimCustomer	Ragged dimension
DimEmployee	Parent-child dimension
DimTime	Role-playing dimension

Question: 49

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) cube.

You need to add a calculated member to the Customer dimension to evaluate the sum of values for the United Kingdom and the United States.

Which expression should you use? (To answer, drag the appropriate expression to the answer area.)

Expressions	Answer Area
[Customer].[Customer Geography].[Country].&[United Kingdom] & [Customer].[Customer Geography].[Country].&[United States]	CREATE MEMBER CURRENTCUBE.[Customer].[Customer Geography].[All].[UK and USA] AS Expression
{[Customer].[Customer Geography].[Country].&[United Kingdom], [Customer].[Customer Geography].[Country].&[United States]}	
[Customer].[Customer Geography].[Country].&[United Kingdom] UNION [Customer].[Customer Geography].[Country].&[United States]	
SUM({[Customer].[Customer Geography].[Country].&[United Kingdom], [Customer].[Customer Geography].[Country].&[United States]})	
SUM(([Customer].[Customer Geography].[Country].&[United Kingdom], [Customer].[Customer Geography].[Country].&[United States]))	

Answer:

```
CREATE MEMBER
CURRENTCUBE.[Customer].[Customer Geography].[All].[UK and USA] AS
```

```
SUM(([Customer].[Customer Geography].[Country].&[United Kingdom], [Customer].[Customer Geography].[Country].&[United States]))
```

Question: 50

You are working with multiple tabular models deployed on a single SQL Server Analysis Services (SSAS) instance.

You need to ascertain the memory consumed by each object in the SSAS instance.

What should you do?

- A. Use the \$System.discover_object_memory_usage dynamic management view.
- B. Use SQL Server Profiler to review session events for active sessions.
- C. Use the Usage Based Optimization wizard to design appropriate aggregations.
- D. Use the Performance Counter group named Processing.

Answer: A

Question: 51

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) multidimensional project that is configured to source data from a SQL Azure database.

You plan to use multiple servers to process different partitions simultaneously. You create and configure a new data source.

You need to create a new partition and configure SQL Server Analysis Services (SSAS) to use a remote server to process data contained within the partition.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Open the cube for editing, select the **Partitions** tab, and then click **New Partition**.

Click **Enable Proactive caching** and then select the **HOLAP** storage mode.

Select the table and then specify the query for the new partition.

Create a linked server for the remote processing location.

Click **Storage Settings** and then click **Options** to open the Storage Options dialog box.

On the Processing and Storage Locations step, ensure that the processing location is set to the Remote Analysis Services data source.

Answer:

Box 1:

Create a linked server for the remote processing location.

Box 2:

Open the cube for editing, select the **Partitions** tab, and then click **New Partition**.

Box 3:

On the Processing and Storage Locations step, ensure that the processing location is set to the Remote Analysis Services data source.

Explanation:

Note:

* You create a remote partition using the Partition Wizard. On the Specify Processing and Storage Options page, for the Remote Analysis Services data source, specify the dedicated database on the remote instance of Analysis Services. This instance of Analysis Services is called the remote server of the remote partition. For Storage location, you can specify the default data location for the remote server or a specified folder on the server.

You must create an Analysis Services database on the remote server and provide appropriate security settings. An additional OLAP data source is created on the remote database pointing to the server on which the partition is defined. The MasterDatasourceID property setting on the remote database points to the data source which, in turn, points to the master server. This property is only set on a database that contains remote partitions. The RemoteDatasourceID property setting on the remote partition specifies the ID of the OLAP data source on the master server that points to the remote server. A remote database can only host remote partitions for a single server.

* Before you create a remote partition, the following conditions must be met:

A database dedicated to serving remote partitions for the local database must be created on the remote server.

The domain user account for the local instance of SQL Server Analysis Services must have administrative access to the remote instance of Analysis Services.

In order to create or maintain a remote partition, your user name must also be included in the OLAP Administrators group for both the remote and the local instances of Analysis Services.

The domain user account for the local instance of Analysis Services must have administrative access to the remote database.

Reference: Creating and Managing a Remote Partition

Question: 52

DRAG DROP

You are planning the installation of PowerPivot for SharePoint.

You install SharePoint Server 2010 Enterprise Edition with Service Pack 1.

You need to install the PowerPivot for SharePoint instance. Then you need to configure the Default Account username used to provision shared services in the SharePoint farm.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Enter the Default Account username and password.

Run the PowerPivot Configuration Tool.

Install SQL Server PowerPivot Add-in for SharePoint.

Use the **Import from PowerPivot** template as the project type.

Open the Services management console and edit the PowerPivot System Service properties.

Open SQL Server Data Tools and create a new project.

Answer:

Box 1:

Install SQL Server PowerPivot Add-in for SharePoint.

Box 2:

Run the PowerPivot Configuration Tool.

Box 3:

Enter the Default Account username and password.

Question: 53

You are developing a SQL Server Analysis Services (SSAS) tabular project.

A model contains tables and columns that must not be visible to the user. The columns and tables cannot be removed because they are used in calculations.

You need to hide the tables and columns.

What should you do?

- A. In the Properties window for the applicable tables and columns, set the Visible property to True.
- B. Right-click the applicable tables and columns and select the Hide option.
- C. Right-click the applicable tables and columns and select the Hide from Client Tools option.
- D. In the Properties window for the applicable tables and columns, set the Enabled property to False.

Answer: C

Question: 54

You are planning to develop a SQL Server Analysis Services (SSAS) tabular project. The project will be deployed to a SSAS server that has 16 GB of RAM.

The project will source data from a SQL Server database that contains a fact table named Sales. The fact table has more than 60 billion rows of data.

You need to select an appropriate design to maximize query performance.

Which data access strategy should you use? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Configure the database to use DirectQuery mode. Create a clustered index which includes all of the foreign key columns of the fact table.
- B. Configure the database to use In-Memory mode. Create a clustered index which includes all of the foreign key columns of the fact table.
- C. Configure the database to use In-Memory mode. Create a columnstore index on all the columns of the fact table.
- D. Configure the database to use DirectQuery mode. Create a columnstore index on all the columns of the fact table.

Answer: D

Question: 55

You are developing a SQL Server Analysis Services (SSAS) tabular project.

You need to grant the minimum permissions necessary to enable users to query data in a tabular model.

Which role permission should you use?

- A. Explorer
- B. Select
- C. Process
- D. Browser
- E. Read Definition
- F. Read

Answer: F

Question: 56

You are developing a SQL Server PowerPivot workbook that sources data from a Microsoft Azure SQL Database database. The PowerPivot model includes a single table named FactSales that consists of four columns named Year, Country, Product and Revenue. The model includes the following two measures.

Sales:=SUM(FactSales[Revenue])

Sales %:=[Sales] / CALCULATE([Sales], ALL(FactSales))

In Microsoft Excel 2010 you create the following PivotTable report.

	A	B	C	D	E
1	Country		Year	2011	
2					
3	Canada		Row Labels	Sales	Sales %
4	Mexico		Bread	234,533	0.54 %
5	USA		Dairy	112,045	0.26 %
6			Meat	534,009	1.22 %
7			Grand Total	880,587	2.01 %

Users report that the Sales % measure computes an incorrect ratio. The measure should meet a requirement to compute a ratio over all visible sales values defined by the query filters. The Grand Total value for the Sales % measure should equal 100%.

You need to fix the Sales % measure to meet the requirement.

Which Data Analysis Expressions (DAX) expression should you use?

- A. = [Sales] / CALCULATE([Sales])
- B. = [sales] / [Sales](ALLSELECTED(FactSales))
- C. = [sales] / CALCULATE([Sales], VALUES(FactSales[Year]), VALUES(FactSales[Country]))
- D. = [sales] / [Sales](ALLEXCEPT(FactSales, FactSales[Year]))

Answer: B

Question: 57

You are developing a SQL Server Analysis Services (SSAS) tabular database. To maximize performance, queries must be resolved only by using cache.

You need to configure the appropriate query mode.

Which query mode should you select?

- A. DirectQuery with In-Memory
- B. In-Memory with DirectQuery

- C. In-Memory
- D. DirectQuery

Answer: C

Question: 58

You are developing a SQL Server Analysis Services (SSAS) tabular project. In the data warehouse, a table named Employee Security defines a relationship between a salesperson's name, logon ID, and assigned sales territory. You need to ensure that each salesperson has access to data from only the sales territory assigned to that salesperson. You also need to minimize the development effort. What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a new role with Read permission and then add each salesperson's logon as a member to the role.
- B. Create a separate tabular project for each sales territory. Grant each salesperson access to the corresponding tabular model of the assigned sales territory.
- C. Add the Employee Security table to the model, define the relationships, and then implement dynamic security by using row filters. Grant each salesperson access to the model.
- D. Create a new Active Directory Domain Services (AD DS) security group and add each salesperson as a member. Then create a new role with Read permission. Add the group as a member to the new role.

Answer: C

Question: 59

You develop a SQL Server Analysis Services (SSAS) tabular project. The tabular model loads data from a SQL Server relational database each day.

You define a connection.

You need to ensure that the connection minimizes the attack surface area of the server.

How should you define the impersonation information for the connection? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Use your domain credentials. Grant least privilege to your account in the source database.
- B. Create and use a new Windows domain account. Grant least privilege to this account in the source database.
- C. Use the credentials of the SQL Server Analysis Services (SSAS) service account. Grant least privilege to this account in the source database.
- D. Use SQL Server authentication.

Answer: B

Question: 60

You are developing a SQL Server Analysis Services (SSAS) tabular project.

A column named City must be added to the table named Customer. The column will be used in the definition of a hierarchy. The City column exists in the Geography table that is related to the Customer table.

You need to add the City column to the Customer table.

How should you write the calculation?

- A. =RELATEDTABLE(Geography)
- B. =RELATED(Geography[City])
- C. =Geography[City]
- D. City:=Geography[City]
- E. City:=RELATED(Geography[City])
- F. City :=RELATEDTABLE (Geography)

Answer: B

Question: 61

You are developing a SQL Server Reporting Services (SSRS) sales summary report.

The report header consists of several images. Report users require PDF exports of the report with no bulky images of the report header.

You need to ensure that the header of the report is hidden when a user exports the report to PDF format.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Set the Hidden property of the report header to (Globals!RenderFormat.Name = "PDF").
- B. Set the Hidden property of the report header to (Globals!RenderFormat.IsInteractive = False).
- C. Set the Hidden property of the report header to FALSE.
- D. Set the Hidden property of the report header to TRUE.

Answer: A

Question: 62

You install SQL Server Reporting Services (SSRS).
You need to back up a copy of the symmetric key.
Which command should you run?

- A. rskeymgmt -a -f %temp%\rs.key -p Password1
- B. rskeymgmt -d
- C. rskeymgmt -i
- D. rskeymgmt -e -f %temp%\rs.key -p Password1

Answer: D

Question: 63

You are developing a new SQL Server Reporting Services (SSRS) report in SQL Server Data Tools (SSDT).
The report contains tables, images, charts, page breaks, gauges, and indicators.
You need to ensure that the reports can be exported to Microsoft Excel and PDF formats and do not have any layout issues.
What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. During development, preview the report in the SSDT preview tab to validate its default HTML rendering.
- B. After development, test the report in a web browser.
- C. Increase the value of the Width property of the report body.
- D. During development, preview the report in the SSDT preview tab and export it to the other rendering formats.

Answer: D

Question: 64

You manage a SQL Server Reporting Services (SSRS) instance in SharePoint integrated mode.
You identify a problem with the operations performed by the Report Server Web service.
You need to view the Report Server trace logs. In which directory should you view the files?

- A. %ProgramFiles%\Microsoft SQL Server\MSRS10_50.MSSQLSERVER\Reporting Services\LogFiles
- B. %ProgramFiles%\Common Files\Microsoft Shared\Web Server Extensions\14\Web Services\ReportServer\LogFiles
- C. %ProgramFiles%\Common Files\ReportServer\LogFiles
- D. %ProgramFiles%\Microsoft SQL Server\MSRS11.MSSQLSERVER\Reporting Services\LogFiles

Answer: D

Question: 65

You are designing a SQL Server Reporting Services (SSRS) report to display vineyard names and their year-to-date (YTD) grape yield.
Grape yield values are classified in three bands:

- High Yield

- Medium Yield
- Low Yield

You add a table to the report. Then you define two columns based on the fields named VineyardName and YTDGrapeYield

You need to set the color of the vineyard text to red, yellow, or blue, depending on the value of the YTD grape yield values.

What should you do?

- A. Use an expression for The Color property of the vineyard text box.
- B. Use an expression for the TextDecoration property of the vineyard text box.
- C. Use an expression for the Style property of the vineyard text box.
- D. Use an expression for the Font property of the vineyard text box.
- E. Add an indicator to the table.

Answer: A

Question: 66

A large manufacturing company has manufacturing plants in many states. Each state has a dedicated SQL Server instance that stores its manufacturing dat

a. Each SQL Server instance is configured identically and all instances have identical database structures.

You are developing a daily report that summarizes information on manufacturing activity.

The report has the following requirements:

- It must have a stale name as one of the report parameters.
- It must provide a daily summary of manufacturing activity of a selected state.
- It must require minimal development and maintenance effort.

You need to develop the report to meet the requirements.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

A. Build one report for each state and instruct users to execute reports as needed.

B. Use a report-specific data source that defines an expression-based connection string based on the state parameter.

C. Build a centralized data repository, schedule a regular Extract Transform, and Load (ETL) process on all manufacturing data, and then use the repository to generate the report.

D. Use a shared data source that defines an expression-based connection string based on the state parameter.

Answer: B

Question: 67

You are creating a new report in SQL Server Report Builder. You add a Microsoft Azure SQL Database data source. Then you add a dataset that has four fields named Year, Country, Category, and Sales.

You must design a matrix as shown in the following table.

	CY 2003	CY2004
Accessories	293,710	407,050
Bikes	9,359,103	9,162,325
Clothing	138,248	201,525
Australia	3,033,784	2,563,884
Canada	535,784	673,628
France	1,026,325	922,179
Germany	1,058,406	1,076,891
United Kingdom	1,298,249	1,210,286
United States	2,838,512	3,324,031
Total	9,791,060	9,770,900

The category rows (the first three rows as shown in the diagram) must present total sales amount by category. The country rows {the next six rows as shown in the diagram} must present total sales amount by country. The total row must present the total sales for each year.

You add a matrix to the report. You add a grouping of the Category field on the rows and a grouping of the Year field on the columns.

You need to add the countries on the rows of the matrix.

Which Row Group option should you select when you add the group?

- A. Adjacent Below
- B. Child Group
- C. Parent Group
- D. Adjacent Above

Answer: A

Question: 68

You are managing a SQL Server Reporting Services (SSRS) instance in native mode. A system role named Developer Support is present on the server.

Members of the Developer Support role cannot modify the report execution timeout period.

You need to enable members of the Developer Support role to modify the report execution timeout period.

Which task should you add to the Developer Support role?

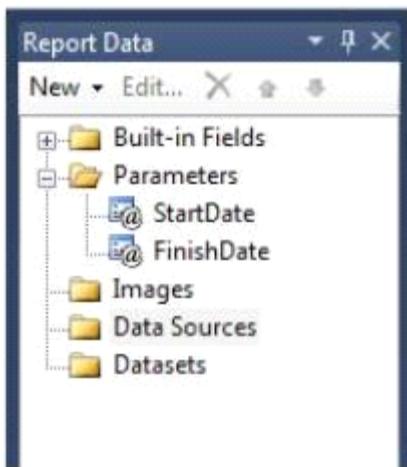
- A. Manage report server properties
- B. Manage shared schedules
- C. Execute report definitions
- D. Manage jobs

Answer: A

Question: 69

DRAG DROP

You are developing a SQL Server Reporting Services (SSRS) report that sources data from a SQL Azure database and a SQL Server Analysis Services (SSAS) cube. The cube contains a date dimension and other dimensions. The report design includes two report parameters named StartDate and FinishDate as shown in the following diagram.



The Data Type property of the parameters is set to Date/Time.

You need to create the dataset based on the SSAS cube. You also need to ensure that the dataset is filtered by the existing report parameters.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Create a data source using the Microsoft SQL Server Analysis Services connection type.

Modify the parameter expressions of the dataset to include the **ToString()** function.

Create two report parameters with Date/Time data types to receive their default values from the two hidden datasets.

Create a dataset with a parameterized filter using a hierarchy from the date dimension that uses the **Contains** operator. Do not close the Dataset Properties window.

Create a dataset with a parameterized filter using a hierarchy from the date dimension that uses the **Range (Inclusive)** operator. Do not close the Dataset Properties window.

Create a data source using the Microsoft SQL Server Analysis Services connection type.

Write an expression in the dataset Parameter Value textbox that converts the Date/Time parameter values to be compatible with the dimension member.

Select the **Parameters** page and then assign an expression to each query parameter to convert the report parameter values to the appropriate date dimension hierarchy member keys.

Answer:

Box 1:

Create a data source using the Microsoft SQL Server Analysis Services connection type.

Box 2:

Create a dataset with a parameterized filter using a hierarchy from the date dimension that uses the **Range (Inclusive)** operator. Do not close the Dataset Properties window.

Box 3:

Select the **Parameters** page and then assign an expression to each query parameter to convert the report parameter values to the appropriate date dimension hierarchy member keys.

Note:

* In Reporting Services, a dataset is based on an existing a data source. A dataset specifies a query, query parameters, filters, and a field collection. You can also specify data options, such as case, collation, kanatype, width, and accent, for the data retrieved from the data source. For more information, see Understanding Report Datasets.

To create a dataset, you must have defined an embedded or shared data source.

* When you deploy Reporting Services, a set of data processing extensions are automatically installed and registered on both the report authoring client and on the report server to provide access to a variety of data source types . SQL Server 2008 Reporting Services (SSRS) installs the following data source types: Microsoft SQL Server, Microsoft SQL Server Analysis Services, Oracle, SAP NetWeaver BI, Hyperion Essbase, Teradata, OLE DB, ODBC, and XML.

* To define a query parameter in MDX in Design mode

In the Report Data pane, right-click on a dataset created from a SQL Server Analysis Services data source type, and then clickQuery. The MDX query designer opens in Design mode.

Drag a dimension to the filter area and drop it on the first cell in the Dimensioncolumn.

In the Hierarchycolumn, choose a value from the drop-down list.

In the Operatorcolumn, choose an operator for the drop-down list.

In the Filter Expressioncolumn, select individual values from the drop-down list, or click the Allmember to choose all values.

In the Parameterscolumn, select the check box to create a report parameter.

Click Run.

After you run the query, click Design on the toolbar to toggle to Query mode to view the MDX query that was created.

Do not change the query text in Query mode if you want to continue to use Design mode to develop the query.

Click Design to toggle back to Design mode.

Click OK.

* To create a dataset

In the Report Data pane, right-click the name of the data source, and then click Add Dataset. The Query page of the Dataset Properties dialog box opens.

In Name, type a name for the dataset or accept the default name.

In Data source, select the name of an existing shared data source, or click New to create a new embedded data source.

Select a Query type option. Options vary depending on the data source type.

Select Text to write a query using the query language of the data source.

Select Table to return all the fields in a relational database table.

Select StoredProcedure to run a stored procedure by name.

In Query, type the query, stored procedure, or table name. Alternatively, click Query Designer to open the graphical or text-based query designer tool, or Import to import the query from an existing report.

In a few cases, the field collection specified by the query can only be determined by running the query on the data source. For example, a stored procedure may return a variable set of fields in the result set. Click Refresh Fields to run the query on the data source and retrieve the field names that are needed to populate the dataset field collection in the Report Data pane. The field collection appears under the dataset node after you close the Dataset Properties dialog box.

In Timeout, type the number of seconds that the report server waits for a response from the database. The default value is 0 seconds. When the time out value is 0 seconds, the query does not time out.

Click OK.

The dataset and its field collection appear in the Report Data pane under the data source node.

Reference: How to: Create a Dataset (Reporting Services); How to: Define Parameters in the MDX Query Designer for Analysis Services (Report Builder 3.0 and SSRS)

Question: 70

You are managing a SQL Server Reporting Services (SSRS) instance.

A website must pass credentials to the local security authority for Reporting Services.

You need to configure Reporting Services to issue a challenge/response when a connection is made without credentials.

Which authentication type should you configure in the RSReportServer.config file?

- A. RSWindowsKerberos only
- B. RSWindowsKerberos and RSWindowsNTLM
- C. RSWindowsBasic
- D. RSWindowsKerberos and RSWindowsNegotiate

Answer: C

Ref: <http://msdn.microsoft.com/en-us/library/ms157273.aspx>

Question: 71

You are creating a SQL Server Analysis Services (SSAS) multidimensional database.

Users need a time dimension for:

- Dates
- Delivery dates
- Ship dates

You need to implement the minimum number of required SSAS objects.

What should you do?

- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure that uses the LastNonEmpty aggregate function. Use a regular relationship between the time dimension and the measure group.
- F. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- G. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- H. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- I. Create several dimensions. Add each dimension to the cube.
- J. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- K. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.

- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: A

Question: 72

You are developing a Microsoft SQL Analysis Services (SSAS) multidimensional project

A fact table named FactHouseSales has a measure column named Are

- a. All values in the column are stored in square feet. Users must be able to analyze the area in different units.
- You create a table named AreaUnit. Each row in the table consists of the unit name and a square feet conversion factor value.
- You need to implement the area conversion in the project.
- What should you do?
- A. Use role playing dimensions.
- B. Use the Business Intelligence Wizard to define dimension intelligence.
- C. Add a measure that uses the Count aggregate function to an existing measure group.
- D. Add a measure that uses the DistinctCount aggregate function to an existing measure group.
- E. Add a measure that uses the LastNonEmpty aggregate function. Use a regular relationship between the time dimension and the measure group.
- F. Add a measure group that has one measure that uses the DistinctCount aggregate function.
- G. Add a calculated measure based on an expression that counts members filtered by the Exists and NonEmpty functions.
- H. Add a hidden measure that uses the Sum aggregate function. Add a calculated measure aggregating the measure along the time dimension.
- I. Create several dimensions. Add each dimension to the cube.
- J. Create a dimension. Then add a cube dimension and link it several times to the measure group.
- K. Create a dimension. Create regular relationships between the cube dimension and the measure group. Configure the relationships to use different dimension attributes.
- L. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a regular relationship between the dimension and measure group.
- M. Create a dimension with one attribute hierarchy. Set the IsAggregatable property to False and then set the DefaultMember property. Use a many-to-many relationship to link the dimension to the measure group.
- N. Create a dimension with one attribute hierarchy. Set the ValueColumn property, set the IsAggregatable property to False, and then set the DefaultMember property. Configure the cube dimension so that it does not have a relationship with the measure group. Add a calculated measure that uses the MemberValue attribute property.
- O. Create a new named calculation in the data source view to calculate a rolling sum. Add a measure that uses the Max aggregate function based on the named calculation.

Answer: N

Question: 73**HOTSPOT**

You are developing a SQL Server Analysis Services (SSAS) cube that contains the data for a running team. The data warehouse used by the cube contains the time durations of laps run by each runner on the team. The time durations are stored in seconds as an integer.

You need to build the following two measures in the cube:

A measure named Measure1 that must contain the average time duration of the laps run by each runner.

A measure named Measure2 that must contain the lap-time duration and the name of the runner who ran the fastest lap.

What should you do? To answer, select the appropriate Aggregation Function property for each measure in the answer area.

Answer Area

Measure1	Measure2
<input type="button" value="▼"/>	<input type="button" value="▼"/>

Answer Area

Measure1	Measure2
<input type="button" value="▼"/> First Child AverageOfChildren Min Max	<input type="button" value="▼"/> First Child AverageOfChildren Min Max

Answer:**Answer Area**

Measure1	Measure2
<input type="button" value="▼"/> First Child AverageOfChildren Min Max	<input type="button" value="▼"/> First Child AverageOfChildren Min Max

Explanation:

Note: For Measure1 we use Average, and for Measure2 Min to get the fastest time.

Question: 74

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) cube. The cube consists of a single measure group.

The measure group consists of one partition that uses MOLAP.

The proactive caching policy has the following requirements:

The cache must be updated when data is changed in the table named **tblOrders**.

Changes must be notified through the use of the XML for Analysis (XMLA) **NotifyTableChange** command.

You need to configure the proactive caching policy to meet the requirements.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

On the **Partitions** tab, click **Storage Settings**.

Select the **SQL Server** notification type, and then select the **tblOrders** table.

Enable proactive caching.

Open the partition storage settings.

Select the **Update the cache periodically** option.

Select the **Client initiated** notification type, and then select the **tblOrders** table.

Answer:

Box 1: On the Partitions tab, Click Storage Settings

We specify the Storage Settings for the correct partition.

Box 2: Enable proactive caching.

We enable proactive caching.

Box 3: Select the Client initiated notification type, and then select the **tblOrders** table.

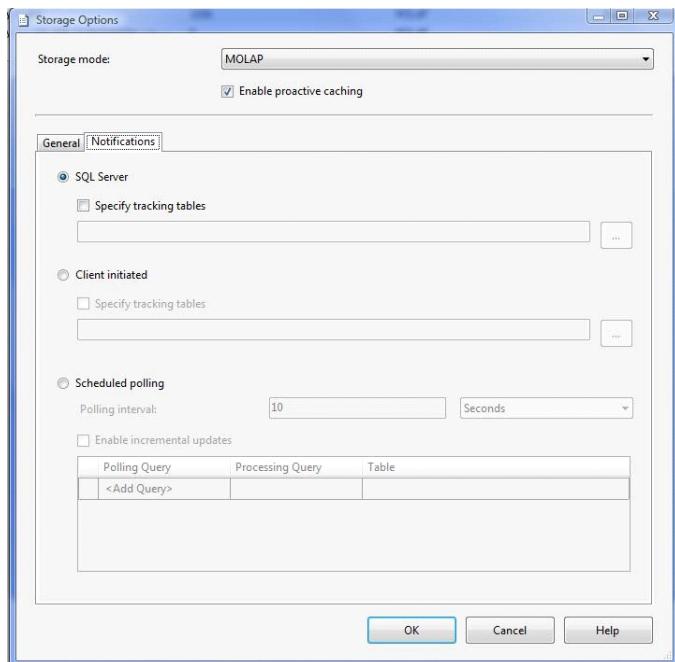
Explanation:

On the Notifications tab, there are three options out of which, as shown below, you can select any one at a time.

* **SQL Server** - With this option, SSAS uses SQL Server notification services/specialized trace mechanism to identify data changes.

* **Client initiated** - With this option, client can specify the XMLA (XML for Analysis) command (**NotifyTableChange**) to identify data changes.

* **Scheduled polling** - With this option, SSAS uses a series of queries to see (polling at defined interval) if there is any data change at the underlying relational database.

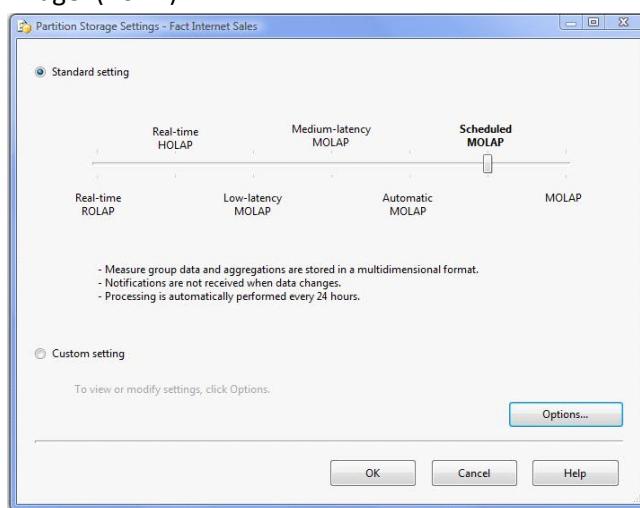
**Note:**

You use the Storage Settings dialog box in BIDS (Business Intelligence Development Studio) to set the proactive caching feature, storage location, and notification settings for a dimension, cube, measure group, or measure group partition.

In the Cube Browser, open your cube and select the Partitions tab.

Expand the measure group and select the partition for which you want to enable proactive caching.

Click the Storage Setting link to open the Partition Storage Settings dialog box similar to the one as shown in below image. (Box 1)



The Custom Setting allows you to explicitly enable proactive caching (if you don't want to use Standard Setting), set storage mode, and notification options. (Box 2)

Question: 75

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) cube.

You need to reuse a Revenue measure group from a different database.

In SQL Server Data Tools (SSDT), which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

From the **Select a Data Source** step, reference the Analysis Services data source.

Launch the Linked Object Wizard.

From the **Select Objects** step, select the Revenue measure group and the dimensions that you need to link.

Launch the Business Intelligence Wizard.

From the **Select Objects** step, select the Revenue measure group that you need to link.

Answer:

Box 1:

Launch the Linked Object Wizard.

Box 2:

From the **Select a Data Source** step, reference the Analysis Services data source.

Box 3:

From the **Select Objects** step, select the Revenue measure group that you need to link.

Explanation:

Note:

- * you can use the Linked Object Wizard to add a measure group from another database
- * You can use the Linked Object Wizard to either link to or import cubes, dimensions, measure groups, calculations, and Key Performance Indicators (KPIs). You can link to or import these items from another database on the same server or from a database on a remote server
- * The Linked Object Wizard guides you through the following steps:

Selecting the Analysis Services data source from which to link or import objects.

Selecting the objects from which to link or import.

Saving the changes.

Create or modify a linked measure

Use SQL Server Data Tools to create a linked measure group.

Finalize any modifications to the original measure group now, in the source cube, so that you don't have to recreate the linked measure groups later in subsequent cubes. You can rename a linked object, but you cannot change any other properties.

In Solution Explorer, double-click the cube to which you are adding the linked measure group. This step opens the cube in Cube Designer.

In Cube Designer, in either the Measures pane or Dimensions pane, right-click anywhere in either pane, then select New Linked Object. This starts the Linked Object Wizard.

On the first page, specify the data source. This step establishes the location of the original measure group. The default is the current cube in the current database, but you can also choose a different Analysis Services database.

On the next page, choose the measure group or dimension you want to link. Dimensions and Cube objects, such as measure groups, are listed separately. Only those objects not already present in the current cube are available.

Click Finish to create the linked object. Linked objects appear in the Measures and Dimensions pane, indicated by the link icon.

Ref: [http://msdn.microsoft.com/en-us/library/ms174899\(v=sql.110\).aspx](http://msdn.microsoft.com/en-us/library/ms174899(v=sql.110).aspx)

Question: 76

You are modifying a SQL Server Analysis Service (SSAS) cube.

The cube consists of a single measure group that contains the following measures:

- Total Quantity On Hand
- Average Quantity On Hand

The measure group has a single partition that uses the MOLAP storage mode.

You need to modify the cube design to ensure that the Total Quantity On Hand measure is updated in real-time and that Average Quantity On Hand measure is updated hourly.

What should you do?

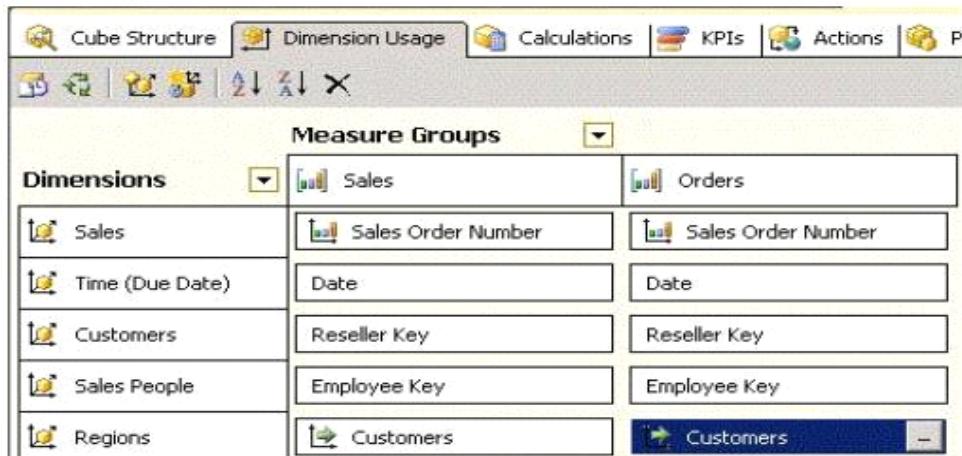
- A. Change the storage mode of the partition to use proactive caching with minimum latency.
- B. Create an XMLA script that will process the cube and then use SQL Server Agent to execute the script continuously.
- C. Create a new measure group for the Average Quantity On Hand measure. Configure the storage mode for the new measure group's partition to ROLAP.
- D. Create a new measure group for the Total Quantity On Hand measure. Configure the storage mode for the new measure group's partition to ROLAP.

Answer: D

Question: 77

You create a materialized reference dimension between a dimension named Regions and a measure group named Orders.

You use a table named Customers as a bridge as shown in the exhibit. (Click the Exhibit button.)



The screenshot shows the 'Dimension Usage' tab of the Analysis Services Dimension Usage tool. The 'Measure Groups' section is selected. A grid displays the relationship between dimensions and measure groups. The 'Dimensions' column lists Sales, Time (Due Date), Customers, Sales People, and Regions. The 'Measure Groups' column lists Sales and Orders. The 'Sales' row under Dimensions has five cells under Measure Groups: Sales Order Number, Sales Order Number, Date, Reseller Key, and Employee Key. The 'Regions' row under Dimensions has one cell under Measure Groups: Customers. The 'Customers' cell in the last row is highlighted with a blue border.

Dimensions	Measure Groups
Sales	Sales Sales Order Number Date Reseller Key Employee Key
Time (Due Date)	Orders Sales Order Number Date
Customers	Reseller Key
Sales People	Employee Key
Regions	Customers

You need to ensure that data is updated to reflect any changes in the relationship. The solution must NOT require that the entire cube be reprocessed.

What should you reprocess?

- A. The Sales partition
- B. The Regions dimension
- C. The Customers dimension
- D. The Orders partition

Answer: C

Question: 78

HOTSPOT

You are designing a dimension named Employee for a SQL Server Analysis Services multidimensional project.

The Employee dimension contains a DateOfBirth attribute and a MaritalStatus attribute.

You need to minimize the amount of time required to process the cube.

What should you do? To answer, select the appropriate relationship type for each attribute in the answer area.

Answer Area

Attribute	Relationship type
DateOfBirth	<input type="text"/>
MaritalStatus	<input type="text"/>

Answer Area

Attribute	Relationship type
DateOfBirth	<input type="text"/> <ul style="list-style-type: none"> Fixed Flexible Rigid
MaritalStatus	<input type="text"/> <ul style="list-style-type: none"> Fixed Flexible Rigid

Answer:

Answer Area

Attribute	Relationship type
DateOfBirth	<input type="text"/> <ul style="list-style-type: none"> Fixed Flexible Rigid
MaritalStatus	<input type="text"/> <ul style="list-style-type: none"> Fixed Flexible Rigid

Question: 79

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) tabular project.

You need to add a calculated column to a table in the model.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

On the Design ribbon inside the Columns group, click **Add**.

Click **Add Formula** and then add a Data Analysis Expressions (DAX) function.

In the formula bar, type an equal sign followed by a Data Analysis Expressions (DAX) expression.

In the formula bar, type an equal sign followed by a Multidimensional Expressions (MDX) expression.

On the **Column** menu, select **Add Column**.

In the model designer, select the table to which you want to add a calculated column.

Answer:

Box 1:

In the model designer, select the table to which you want to add a calculated column.

Box 2:

On the **Column** menu, select **Add Column**.

Box 3:

In the formula bar, type an equal sign followed by a Data Analysis Expressions (DAX) expression.

Explanation:

Note:

* To create a new calculated column

In the model designer, in Data View, select the table to which you want to add a calculated column, then click the Column menu, and then click Add Column.

Add Column is highlighted over the empty rightmost column, and the cursor moves to the formula bar.

To create a new column between two existing columns, right-click an existing column, and then click Insert Column.

In the formula bar, do one of the following:

Type an equal sign followed by a formula.

Type an equal sign, followed by a DAX function, followed by arguments and parameters as required by the function.

Click the function button (fx), then in the Insert Function dialog box, select a category and function, and then click OK.

In the formula bar, type the remaining arguments and parameters as required by the function.

Press ENTER to accept the formula.

* Calculated columns, in tabular models, allow you to add new data to your model. Instead of pasting or importing values into the column, you create a DAX formula that defines the column's row level values. The calculated column can then be used in a report, PivotTable, or PivotChart as would any other column.

* A calculated column is DAX expression that creates a new column in a table and the obtained values are stored in the table; the calculated column expression is evaluated every time the table is processed.

* In tabular object models the calculated column is a column in a table whose values are calculated upon definition of the column, from an expression.

Reference: Create a Calculated Column (SSAS Tabular)

Question: 80

You are developing a BI Semantic Model (BISM) based on a simple and small dataset sourced from SQL Server. The data size and complexity of the data relationships will not change. The model will be used to produce reports in Power View. The reports will show the relationship between product sales and rainfall over time.

You need to use an appropriate project type.

Which project types should you use? (Each answer presents a complete solution. Choose all that apply.)

- A. A tabular project that uses the In-Memory query mode
- B. A tabular project that uses the DirectQuery query mode
- C. A multidimensional project that uses the MOLAP storage mode and proactive cache
- D. A multidimensional project that uses the ROLAP storage mode and columnstore indexes
- E. A PowerPivot workbook that is deployed to Microsoft SharePoint Server 2010

Answer: A, B, E

Question: 81

You work in the Business Intelligence (BI) department of a multinational company.

To share its sales data between the various subsidiaries, the company has requested a new corporate BI solution that meets the following requirements:

The solution must use SQL Server Analysis Services (SSAS) multidimensional or tabular Business Intelligence Semantic Model (BISM).

The model must incrementally add 10 million fact rows of sales data per month.

The model must be translated to English, German, Chinese, or Spanish based on users' locale.

The model must be able to contain the most recent 36 months of data, in order to let users query the data.

You need to select the appropriate model type and partitioning strategy to meet the requirements.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create and deploy a BISM multidimensional model with one partition for all of the data.
- B. Create and deploy a BISM tabular model with one partition for each of the 36 months.
- C. Create and deploy a BISM multidimensional model with one partition for each of the 36 months.
- D. Create and deploy a BISM tabular model with one partition for all of the data.

Answer: C

Question: 82

A multidimensional SQL Server Analysis Services (SSAS) database will be tested next week.

During the test period, users will access the database for 30 days. Multidimensional Expressions (MDX) queries generated during the test period must represent the variety of queries that will be used in the production environment.

After testing completes, you need to implement aggregations for every partition in the solution while minimizing development effort. You need to ensure that the aggregations are optimal.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Set up a query log and record all user queries during the test period. After completion of the test, use the Usage Based Optimization Wizard to define aggregations for each partition.
- B. During the test period, identify long-running queries by using SQL Server Profiler. Use those queries to design aggregations by using the DesignAggregations command in XML for Analysis (XMLA).
- C. During the test period, run SQL Server Profiler for 10 minutes every day and record all queries executed in those 10 minutes. Use the Aggregation Design Wizard to design the aggregations.
- D. Develop a SQL Server Integration Services (SSIS) package by using a Script task and then use Analysis Management Objects (AMO) to design the aggregations.

Answer: D

Question: 83**HOTSPOT**

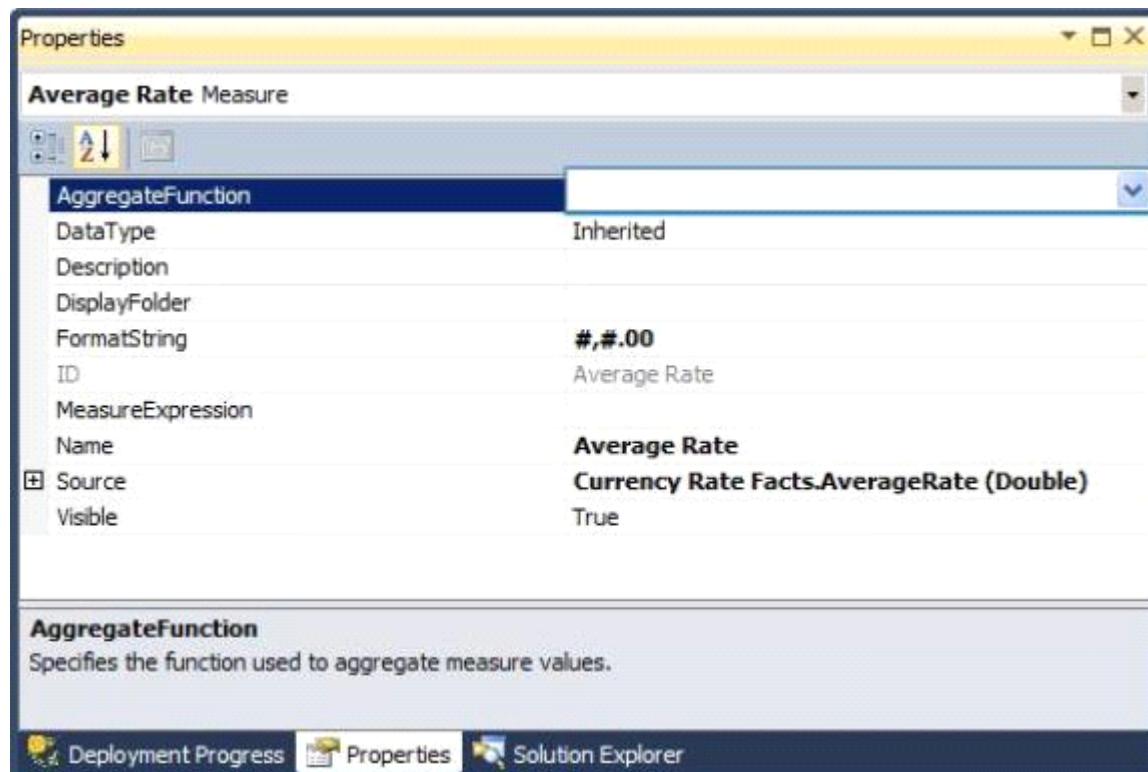
You are developing a SQL Server Analysis Services (SSAS) cube for the accounts department.

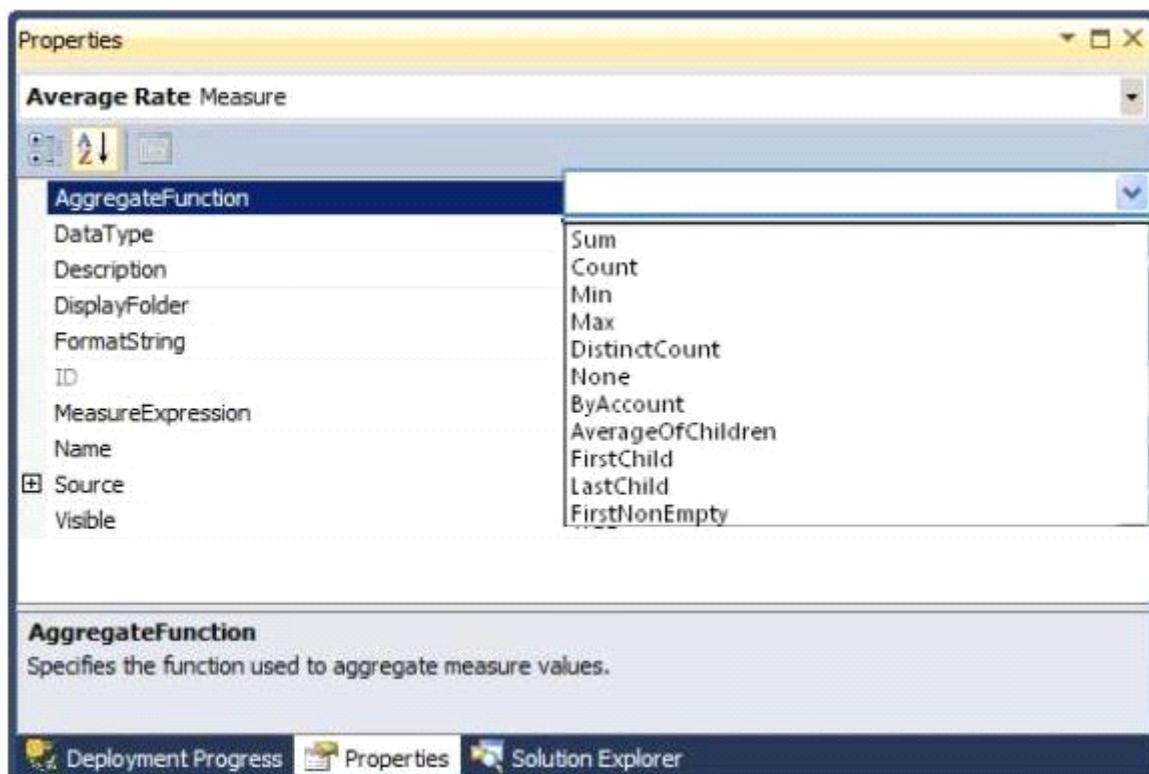
You create a measure group named Exchange Rate that consists of measures pertaining to currency exchange rates. One of the measures in this group is named Average Rate and it will be used to report the average currency exchange rate over time.

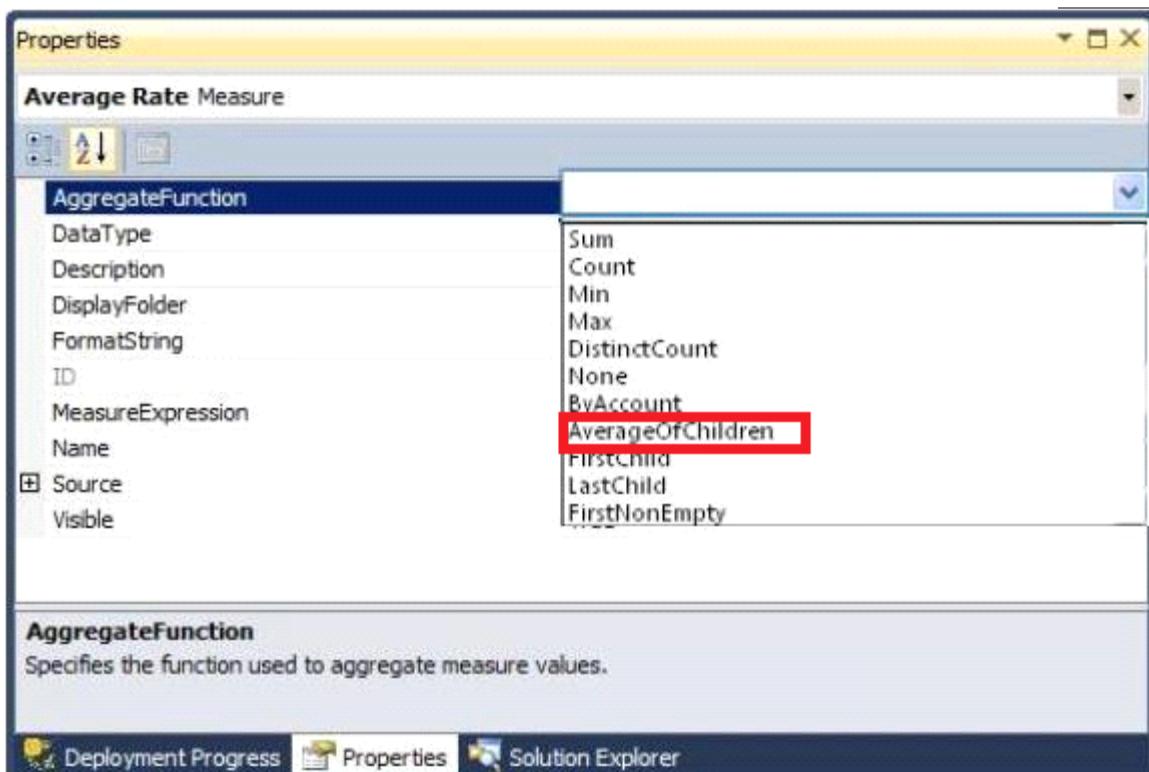
Currently the AggregationFunction property for the Average Rate measure is set to Sum.

You need to ensure that Average Rate measure reports the average of the currency exchange rate over time.

Which value should you select for the AggregationFunction property for the Average Rate measure? To answer, select the appropriate setting in the answer area.





Answer:

Explanation:

AverageOfChildren

Specifies average of leaf descendants in time. Average does not count an empty value as 0.

Question: 84

HOTSPOT

You are developing a SQL Server Analysis Services (SSAS) cube.

Revenue must be compared to a goal and described by a status and a trend. Revenue, goal, status, and trend will be defined by Multidimensional Expressions (MDX) expressions.

You need to add the Revenue measure.

Which tab should you select? (To answer, select the appropriate tab in the answer area.)

Answer Area



Answer:



Question: 85

You are designing a SQL Server Analysis Services (SSAS) cube based on a Microsoft Azure SQL Database data warehouse.

You need to implement a degenerate dimension.

What should you do?

- A. Use the fact table as the data source for the dimension.
- B. Create snowflake dimension tables based on normalized views of the fact table in the data source.
- C. Create a junk dimension table based on the fact table in the data source.
- D. Add a surrogate key to the fact table and use it as the degenerate dimension key.

Answer: A

Question: 86

You are developing a SQL Server Analysis Services (SSAS) cube. The cube contains several dimensions, a local measure group, and a linked measure group. Both measure groups use MOLAP partitions.

You need to write-enable one of the linked measure group partitions to support Microsoft Excel 2010 PivotTable What-If Analysis.

What should you do before the partition can be write-enabled?

- A. Implement the cube as a local cube.
- B. Ensure that the measure group measures only use semiadditive aggregation functions.
- C. Implement the linked measure group as a local measure group.
- D. Ensure that the measure group measures only use nonadditive aggregation functions

Answer: C

Question: 87

You are designing a SQL Server Analysis Services (SSAS) cube for the sales department at your company.

The sales department has the following requirements for the cube:

Include a year-over-year (YOY) calculation

Include a month-over-month (MOM) calculation

You need to ensure that the calculations are implemented in the cube.

Which Multidimensional Expressions (MDX) function should you use?

- A. UNREGINTERCEPT()
- B. LASTPERIODS()
- C. TIMEINTELLIGENCE()
- D. PARALLELPERIOD()

Answer: D

Question: 88

DRAG DROP

You are using Multidimensional Expressions (MDX) to query a SQL Server Analysis Services (SSAS) cube.

You need to compute the aggregate value of the 10 most-ordered products in the Product Categories hierarchy. The Product level is the lowest in the hierarchy.

Which functions should you use to complete the MDX query? (To answer, drag the appropriate functions from the list of functions to the correct locations in the answer area.)

LEAVES	WITH MEMBER [Measures].[SumOfTop10products]
TOPSUM	AS (
MEMBERS	(
TOPCOUNT	[Product].[Product Categories], ,)
AGGREGATE	, 10
DESCENDANT	,[Measures].[Order Quantity])

```

    , [Measures].[Order Quantity])
    , [Measures].[Order Quantity])

SELECT {[Measures].[Order Quantity], [Measures].[SumOfTop10products]}
ON COLUMNS
, {[Product].[Product Categories].[Category].&[4]
, [Product].[Product Categories].[Subcategory].&[31]}
ON ROWS
FROM [Orders]

```

Answer:

```

WITH MEMBER [Measures].[SumOfTop10products]
AS TOPCOUNT (
    AGGREGATE (
        DESCENDANT ([Product].[Product Categories]), , LEAVES
        , 10
        , [Measures].[Order Quantity])
        , [Measures].[Order Quantity])
SELECT {[Measures].[Order Quantity], [Measures].[SumOfTop10products]}
ON COLUMNS
, {[Product].[Product Categories].[Category].&[4]
,[Product].[Product Categories].[Subcategory].&[31]}
ON ROWS
FROM [Orders]

```

Explanation:

Note:

* Example (order of TopCount and Aggregate):

WITH

SET

[Top25Customers] as

TopCount([Customers].[All Customers].Children, 25.0, [Measures].[Sales])

MEMBER [Customers].[All Customers].[Rest of Customers] as

Aggregate(Except([Customers].[All Customers].Children,[Top25Customers]))

SELECT

NON EMPTY {CROSSJOIN([Markets].[All Markets].Children,{[Measures].Sales})}

ON COLUMNS,

Union([Top25Customers],{[Customers].[All Customers].[Rest of Customers]}) ON ROWS

from [SteelWheelsSales]

* TopCount: Sorts a set in descending order and returns the specified number of elements with the highest values.

* Aggregate:

Returns a number that is calculated by aggregating over the cells returned by the set expression. If a numeric expression is not provided, this function aggregates each measure within the current query context by using the default aggregation operator that is specified for each measure. If a numeric expression is provided, this function first evaluates, and then sums, the numeric expression for each cell in the specified set.

* Example:

One can extract the leaf members of a parent child hierarchy by asking the descendants of the root member with the following expression:

Descendants([Organization].[Organizations].&[1], , LEAVES)

* Incorrect:

/ Not TopSUM: Returns, in order of decreasing rank, the top-most rows of a table whose cumulative total is at least a specified value.

Question: 89

You are restructuring an existing cube. One of the measures in the cube is Amount. The Sum aggregation function is used for the Amount measure. The cube includes a dimension named Account and the dimension's Type property is set to Accounts. The Account dimension includes an account type attribute.

You need to ensure that the Amount measure aggregates correctly according to the account type classification. Development effort must be minimized.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Develop a .NET application that uses Analysis Management Objects (AMO) to change the existing AggregateFunction property value of the Amount measure to FirstNonEmpty and then use the application.
- B. Develop a .NET application that uses Analysis Management Objects (AMO) to change the existing AggregateFunction property value of the Amount measure to ByAccount and then use the application.
- C. Use SQL Server Data Tools to change the AggregateFunction property value of the Amount measure to ByAccount.
- D. Add the ByAccount attribute to the account dimension.

Answer: C

Question: 90

You are developing a SQL Server Analysis Services (SSAS) cube.

You must create a four-level hierarchy for the employee dimension. Each level must be associated with an attribute in the employee dimension table. Two thirds of the dimension data contain values for all four attributes. The remainder of the dimension data contains values for the first three of the four attributes only.

You need to create the hierarchy so that logically missing members will not be shown by the reporting tool.

Which type of hierarchy should you create?

- A. A parent-child hierarchy
- B. A sparse hierarchy
- C. A ragged hierarchy
- D. A balanced hierarchy

Answer: C

Question: 91

You are adding dimensions and a new measure group named Accounts Receivable to an existing SQL Server Analysis Services (SSAS) cube.

Date is one of the existing dimensions in the SSAS database. The underlying fact table for the measure group is associated with multiple dates, including InvoiceDate, DueDate and PaymentDate.

You need to ensure that users can slice the Accounts Receivable measures by InvoiceDate, DueDate, and PaymentDate. You also need to ensure that the time required to process the database is minimized.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create three new perspectives named InvoiceDate, DueDate, and PaymentDate.
- B. Create three independent dimensions named InvoiceDate, DueDate and PaymentDate and then link all three of them to the Accounts Receivable measure group.
- C. Create cube dimensions named InvoiceDate, DueDate, and PaymentDate by using the existing Date dimension in the database.
- D. Add three attributes named InvoiceDate, DueDate, and PaymentDate to the existing Date dimension in the database.

Answer: C

Question: 92

You are developing a SQL Server Analysis Services (SSAS) cube named Sales Planning. The cube consists of two measure groups named Sales and Planning. Each measure group is based on a data warehouse fact table and consists of a single MOLAP partition that has the same name as its measure group.

The Planning measure group consists of two measures:

- Forecast, which uses the Sum aggregate function
- Forecast Count, which uses the Count aggregate function

Users contribute planning values by using a legacy application. An extract, transform, load (ETL) process is scheduled to periodically transfer the planning values from the database of the legacy application to the data warehouse.

Financial analysts query the Sales Planning cube and report that the planning values are sometimes out of date. A

new company requirement mandates that the planning values be entered directly into the cube by using Microsoft Excel 2010 PivotTable What-If Analysis.

You need to write-enable the Planning partition.

What should you do before write-enabling the partition?

- A. Set the StorageMode property of the Planning partition to Rolap
- B. Set the ProcessingMode property of the Planning partition to LazyAggregations.
- C. Set the ProcessingMode property of the Planning measure group to LazyAggregations.
- D. Remove the Forecast Count measure.
- E. Set the Type property of the Planning measure group to Budget.
- F. Convert the Planning measure group to a linked measure group.

Answer: D

Explanation:

A cube can be write-enabled only if all its measures use the Sum aggregate function.

Question: 93

You are developing a BI Semantic Model (BISM) that retrieves data from several sources including a Microsoft Azure SQL Database database and an OData data feed. The model will be deployed to a server with significantly more memory than the total size of the source data.

You have the data feed URL, which you will use when developing the model in SQL Server Data Tools (SSDT).

The model must meet the following requirements:

- Maximize performance
- Data latency of up to one month is acceptable

You need to choose a project type and a data access mode to meet the requirements.

What should you do?

- A. Select the multidimensional project type and use the ROLAP storage mode.
- B. Select the tabular project type and use the In-Memory query mode.
- C. Select the tabular project type and use the DirectQuery query mode.
- D. Select the multidimensional project type and use the MOLAP storage mode.

Answer: B

Question: 94

You are modifying a SQL Server Analysis Services (SSAS) cube.

The cube consists of a single measure group that contains the following measures:

- Total Quantity On Hand
- Average Quantity On Hand

The measure group has a single partition that uses the MOLAP storage mode.

You need to modify the cube design to ensure that the Total Quantity On Hand measure is updated in real-time and that Average Quantity On Hand measure is updated hourly.

What should you do?

- A. Create a new measure group for the Total Quantity On Hand measure. Configure the storage mode for the new measure group's partition to ROLAP.
- B. Create a drillthrough action that will query the underlying data source in real time for the Total Quantity On Hand measure.
- C. Change the storage mode of the partition to ROLAP.
- D. Add an additional MOLAP partition to the measure group.

Answer: A**Question: 95**

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) cube.

You need to add a calculated member to the Customer dimension to evaluate the sum of values for France and Germany.

Which expression should you use? (To answer, drag the appropriate expression to the answer area.)

Expressions	Answer Area
[Customer].[Customer Geography].[Country].&[France] & [Customer].[Customer Geography].[Country].&[Germany]	CREATE MEMBER CURRENTCUBE.[Customer].[Customer Geography].[All].[Average FR and DE] AS
{[Customer].[Customer Geography].[Country].&[France], [Customer].[Customer Geography].[Country].&[Germany]}	Expression
[Customer].[Customer Geography].[Country].&[France] UNION [Customer].[Customer Geography].[Country].&[Germany]	***
SUM({[Customer].[Customer Geography].[Country].&[France], [Customer].[Customer Geography].[Country].&[Germany]})	
SUM(([Customer].[Customer Geography].[Country].&[France], [Customer].[Customer Geography].[Country].&[Germany]))	

Answer:

Answer Area

```
CREATE MEMBER  
CURRENTCUBE.[Customer].[Customer Geography].[All].[Average FR and DE] AS  
  
SUM([Customer].[Customer Geography].[Country].[  
[France], [Customer].[Customer Geography].[Country].[  
[Germany]])
```

Question: 96

You are conducting a design review of a multidimensional project.

In the Customer Geography dimension, all non-key attributes relate directly to the key attribute.

The underlying data of the Customer Geography dimension supports relationships between attributes.

You need to increase query and dimension processing performance.

What should you do?

- A. For the dimension attributes of the Customer Geography dimension, define appropriate attribute relationships.
- B. For the dimension attributes of the Customer Geography dimension, set the GroupingBehavior property to EncourageGrouping.
- C. For the Customer Geography dimension, set the Processing Mode property to LazyAggregations.
- D. For the Customer Geography dimension, set the ProcessingPriority property to 1.

Answer: A

Question: 97

You are modifying a SQL Server Analysis Services (SSAS) cube that aggregates order data from a Microsoft Azure SQL Database database. The existing database contains a customer dimension.

The marketing team has requested that customer marketing categories be added to the database.

The marketing categories must meet the following requirements:

A customer member must be able to belong to multiple category members.

A category member must be able to group several customer members.

The marketing team must be able to create new categories every month in the data source.

You need to implement the appropriate solution to meet the requirements while ensuring that the amount of development and maintenance time is minimized.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a dimension named Marketing Category Name and then configure a many-to-many relationship.
- B. Create a dimension named Marketing Category Name and then configure a regular relationship.
- C. Add an attribute hierarchy named Marketing Category Name to the customer dimension.
- D. Add an attribute hierarchy for each marketing category to the customer dimension. Configure each hierarchy to have two members named Yes and No.

Answer: A

Question: 98

DRAG DROP

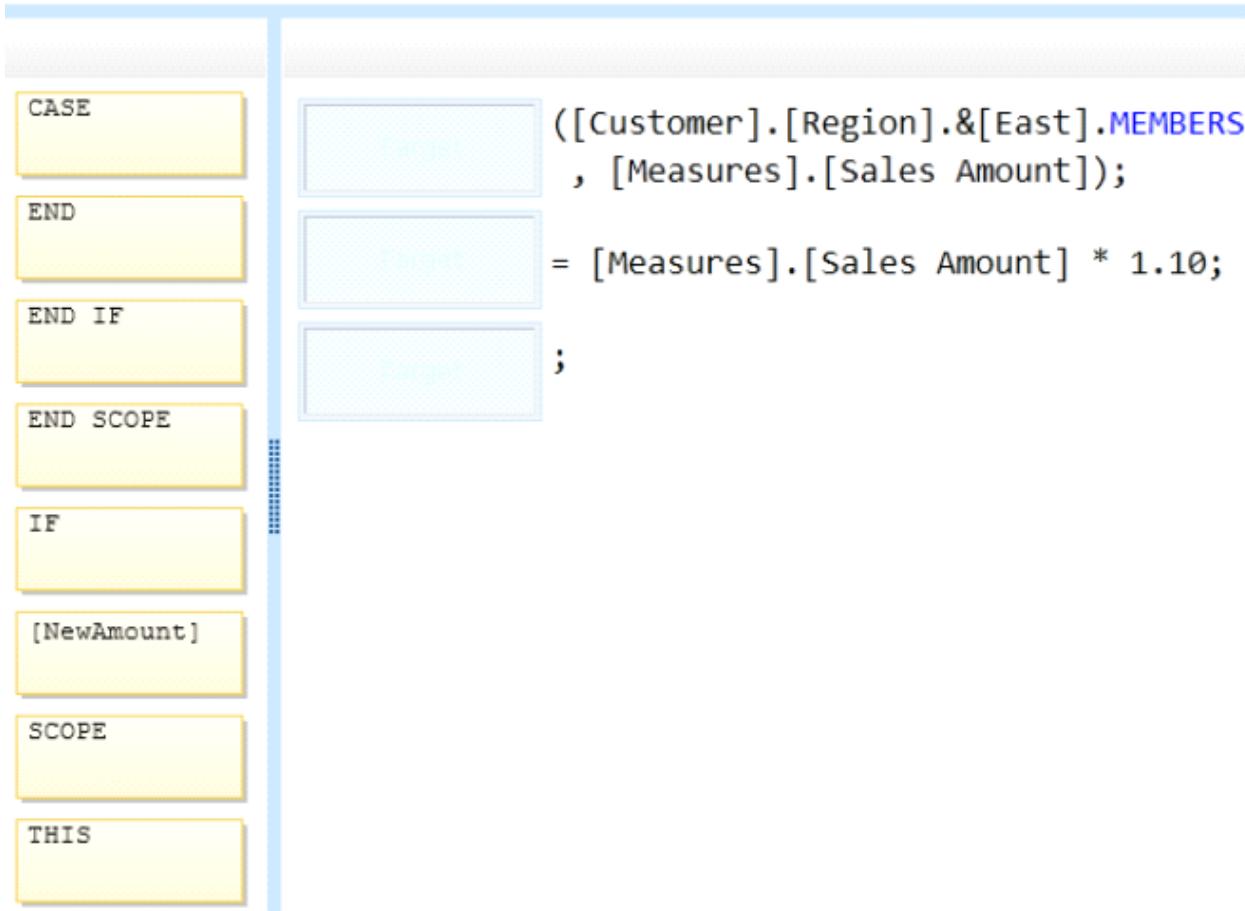
You are making changes to a cube named Sales.

You must increase the value of the measure named Sales Amount by 10%. The increase must be applied only to the children of the member named East in the dimension named Customer.

You need to complete the Multidimensional Expressions (MDX) statement in the calculations section of the Sales cube.

Which statement fragments should you use? (To answer, drag the appropriate statement component to the correct location or locations in the answer area)

- a. Use only components that apply.)



Answer:

```
SCOPE ([Customer].[Region].&[East].MEMBERS
      , [Measures].[Sales Amount]);
THIS = [Measures].[Sales Amount] * 1.10;
END SCOPE ;
```

Explanation:

Note:

SCOPE THIS END SCOPE

* SCOPE Statement (MDX)

Limits the scope of specified Multidimensional Expressions (MDX) statements to a specified subcube.

* Example:

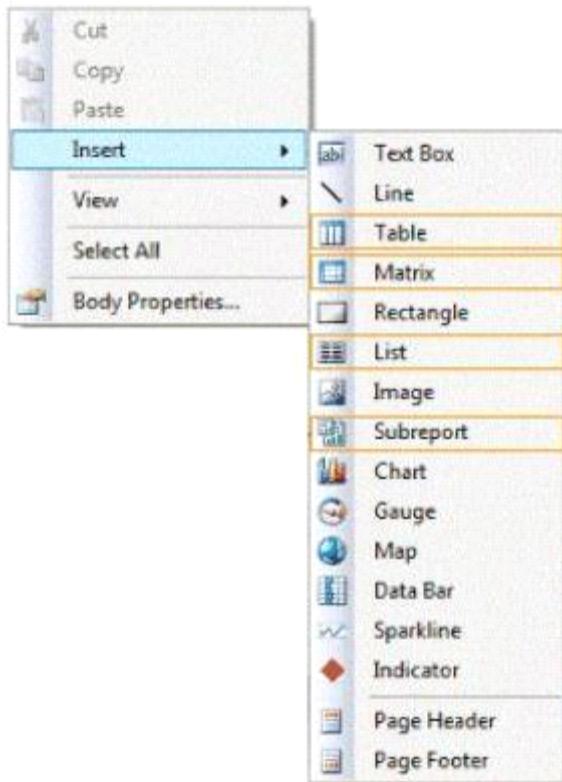
```
cope
(
    [Date].[Fiscal Year].&[2002],
    [Date].[Fiscal].[Month].Members
);
This = [Date].[Fiscal].CurrentMember.Parent / 3 ;
End Scope;
```

Question: 99

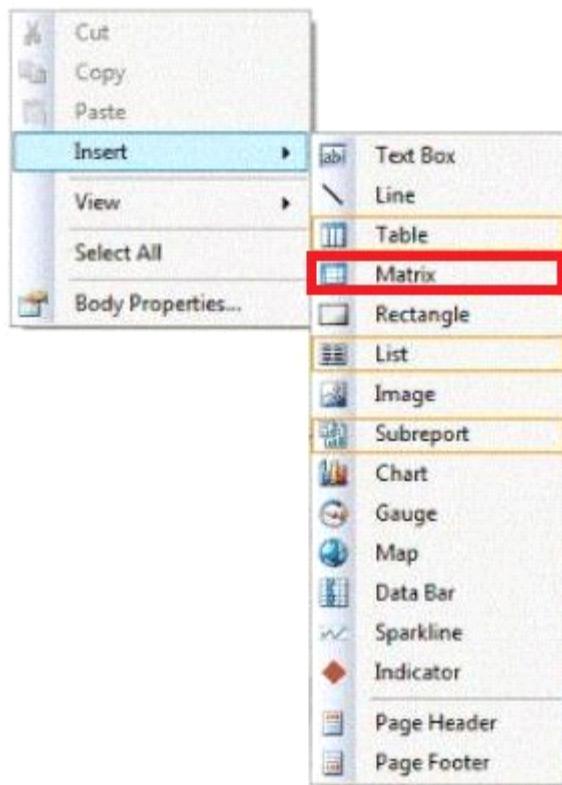
HOTSPOT

You are designing a SQL Server Reporting Services (SSRS) report that sources data from a SQL Azure database. You need to design the report to show the sum of sales. The sales must be grouped by region on the rows and year on the columns.

Which report item should you add? To answer, select the appropriate setting in the answer area.



Answer:



Question: 100

You are working with a SQL Server Reporting Services (SSRS) instance in native mode. An item role named Developer is present on the server.

The Developer role cannot view and modify report caching parameters.

You need to ensure that the Developer role can view and modify report caching parameters.

Which task should you add to the Developer role?

- A. Manage individual subscriptions
- B. View data sources
- C. Manage report history
- D. Manage all subscriptions

Answer: C

Question: 101

You manage a SQL Server Reporting Services (SSRS) instance. The ReportingServicesService.exe.config file has been modified to enable logging.

Some users report that they cannot access the server.

You need to ascertain the IP addresses of the client computers that are accessing the server.

What should you do?

- A. View the ExecutionLog view in the Report Server database.
- B. View the Report Server service trace log.
- C. View the Report Server HTTP log.

D. View the Windows System event log.

Answer: C

Question: 102

A multinational retailer has retail locations on several continents. A single SQL Server Reporting Services (SSRS) instance is used for global reporting.

A SQL Server Analysis Services (SSAS) instance for each continent hosts a multidimensional database named RetailSales. Each RetailSales database stores data only for the continent in which it resides. All of the SSAS instances are configured identically. The cube names and objects are identical.

Reports must meet the following requirements:

A report parameter named ServerName must be defined in each report.

When running a report, users must be prompted to select a server instance.

The report data source must use the Microsoft SQL Server Analysis Services data source type.

You need to create a data source to meet the requirements.

How should you define the expression that is assigned to the connection string property of the data source?

A.="Server=" & Parameters!ServerName.Value & ";Initial Catalog=RetailSales"

B.="Data Source=@ServerName; Initial Catalog=RetailSales"

C.="Data Source=" & Parameters!ServerName.Value & ";Initial Catalog=RetailSales"

D.="Server=" & Parameters!ServerName.Value

E.="Server=@ServerName; Initial Catalog=RetailSales"

Answer: C

Question: 103

You are designing a SQL Server Reporting Services (SSRS) report based on a SQL Server Analysis Services (SSAS) cube. The cube is used to measure sales growth by salesperson.

The cube contains a Key Performance Indicator (KPI) to show if a salesperson's sales are off target slightly off target, or on target.

You need to add a report item that visually displays the KPI status value as a red, yellow, or green flag.

Which report item should you add?

A. An Indicator

B. A Gauge that uses the Radial type

C. A Gauge that uses the Linear type

D. A Sparkline

E. A Data Bar

Answer: A

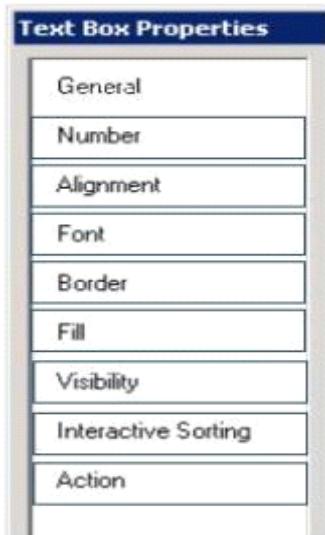
Question: 104

HOTSPOT

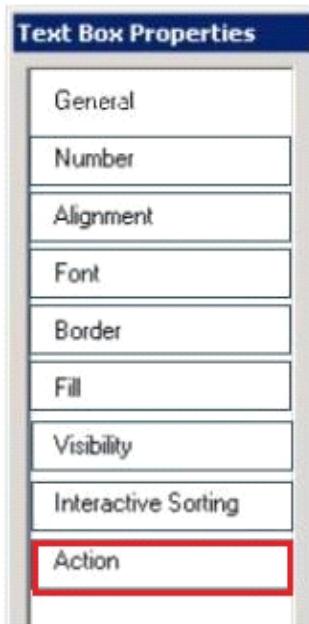
You create a new report in SQL Server Data Tools (SSDT). The report queries a Microsoft Azure SQL Database database table.

In a report table showing sales by cities, you need to enable users to interactively drill through to another report to show details of the customer sales within a chosen city.

Which text box property page should you use? (To answer, select the appropriate tab in the answer area.)



Answer:



Question: 105

You are developing a SQL Server Reporting Services (SSRS) report that renders in HTML. The report includes a dataset with fields named Description, Price, and Color. The report layout includes a table that displays product details and also includes columns named Description, Price, and Color.

You need to modify the report so that users can sort products by the Price column.
What should you do?

- A. Set the SortExpression value to = Fields!Price.Description for the Price text box.
- B. Set the SortExpression property to =Fields!Price.Name for the Price text box.
- C. Set the SortExpression property to =Fields!Price.Value for the Price text box.
- D. Add a custom action to the Price text box.

Answer: C

Question: 106

You are developing a SQL Server Analysis Services (SSAS) tabular project.

A model contains tables and columns that must not be visible to the user. The columns and tables cannot be removed because they are used in calculations. The calculations are used to calculate the budget and forecast for the current quarter.

You need to hide the tables and columns.
What should you do?

- A. Before adding the forecast calculations to the model, right-click the applicable tables and columns and select the Hide option.
- B. After adding the budget calculations to the model, in the Properties window for the applicable tables and columns, set the Visible property to True.
- C. Before adding the forecast calculations to the model, right-click the applicable tables and columns and select the Hide from Client Tools option.
- D. After adding the budget calculations to the model, in the Properties window for the applicable tables and columns, set the Enabled property to False.

Answer: C

Question: 107

You are developing a SQL Server Analysis Services (SSAS) tabular project. A model defines a measure named Revenue and includes a table named Date. The table includes year, semester, quarter, month, and date columns. The Date column is of data type Date. The table contains a set of contiguous dates.

You need to create a measure to report on year-over-year growth of revenue.
What should you do? (Each answer presents a complete solution. Choose all that apply.)

- A. Define the following calculation.

```
Year Over Year Revenue Growth:=[Revenue] - CALCULATE([Revenue],  
SAMEPERIODLASTYEAR('Date'[Date]))
```

- B. Define the following calculation.

```
Year Over Year Revenue Growth:=CALCULATE([Revenue], DATEADD('Date'[Date], 1, YEAR))
```

- C. Use the Business Intelligence Wizard and then use the **Define time intelligence** enhancement.

- D. Define the following calculation.

```
Year Over Year Revenue Growth:=[Revenue] - CALCULATE([Revenue],  
PARALLELPERIOD('Date'[Date], -12, MONTH))
```

- A. Option A
- B. Option B
- C. Option C
- D. Option D

Answer: A, C, D

Question: 108

You are developing a SQL Server Analysis Services (SSAS) tabular project. The model has tables named Invoice Line Items and Products.

The Invoice Line Items table has the following columns:

Product Id

Unit Sales Price

The Unit Sales Price column stores the unit price of the product sold.

The Products table has the following columns:

Product Id

Maximum Sales Price

The Maximum Sales Price column is available only in the Products table.

You add a column named Is Overpriced to the Invoice Line Items table. The Is Overpriced column must store a value of TRUE if the value of the Unit Sales Price is greater than the value of the Maximum Sales Price. Otherwise, a value of FALSE must be stored.

You need to define the Data Analysis Expressions (DAX) expression for the Is Overpriced column.

Which DAX formula should you use? (Each answer represents a complete solution. Choose all that apply.)

- A. `=IF(LOOKUPVALUE(Products[Unit Sales Price], Products[Product Id], [Product Id]) > [Maximum Sales Price]), TRUE, FALSE)`
 - B. `=IF([Unit Sales Price] > RELATED(Products[Maximum Sales Price])), TRUE, FALSE)`
 - C. `=IF([Unit Sales Price] > LOOKUPVALUE(Products[Maximum Sales Price], Products[Product Id], [Product Id]), TRUE, FALSE)`
 - D. `=IF(RELATED(Products[Unit Sales Price])) > [Maximum Sales Price], TRUE, FALSE)`
- A. Option A
B. Option B
C. Option C
D. Option D

Answer: B, C

Question: 109

You are developing a SQL Server Analysis Services (SSAS) tabular project. The model includes a table named DimEmployee. The table contains employee details, including the sales territory for each employee. The table also defines a column named EmployeeAlias which contains the Active Directory Domain Services (AD DS) domain and logon name for each employee. You create a role named Employees. You need to configure the Employees roles so that users can query only sales orders for their respective sales territory. What should you do?

- A. Add a row filter that implements the LOOKUPVALUE and USERNAME functions.
- B. Add a row filter that implements only the CUSTOMDATA function.
- C. Add a row filter that implements the LOOKUPVALUE and CUSTOMDATA functions.
- D. Add a row filter that implements only the USERNAME function.

Answer: A

Question: 110

You are developing a SQL Server Analysis Services (SSAS) tabular project for a Power View solution. You need to grant permission for salespersons to view only the data based on their sales territory. What should you do?

- A. Create a member and then create a Data Analysis Expressions (DAX) filter.
- B. Create a member and then create a Multidimensional Expressions (MDX) filter.
- C. Use SQL Server Management Studio to create a role. Then create a Data Analysis Expressions (DAX) filter.
- D. Use SQL Server Management Studio to create a role. Then create a Multidimensional Expressions (MDX) filter.

Answer: C

Question: 111

You are developing a SQL Server Analysis Services (SSAS) tabular database.

To maximize performance, the queries must be resolved by using cache unless otherwise specified in the connection string.

You need to configure the appropriate query mode.

Which query mode should you select?

- A. In-Memory with DirectQuery
- B. DirectQuery with In-Memory
- C. In-Memory
- D. DirectQuery

Answer: A

Question: 112

You have a tabular model hosted in SQL Server Analysis Services (SSAS).

You need to add new rows to an existing table. The solution must load only the new rows to the table.

Which processing mode should you use?

- A. Process Data
- B. Process Full
- C. Process Add
- D. Process Clear
- E. Process Default

Answer: C

Question: 113

DRAG DROP

You have a single SQL Server 2008 R2 Analysis Services (SSAS) instance. You are planning to upgrade the instance to SQL Server 2014.

You need to import an existing PowerPivot workbook to create a tabular project.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Install SQL Server PowerPivot Add-in for SharePoint.

Install an instance of SQL Server 2012 Analysis Services in Multidimensional and Data Mining mode.

Use the **Import from Server (Tabular)** template as the project type.

Use the **Import from PowerPivot** template as the project type.

Open SQL Server Data Tools and create a new project.

Install an instance of SQL Server 2012 Analysis Services in Tabular mode.

Open SQL Server Business Intelligence Development Studio (BIDS) and create a new project.

Answer:

Box 1:

Install an instance of SQL Server 2012 Analysis Services in Tabular mode.

Box 2:

Open SQL Server Data Tools and create a new project.

Box 3:

Use the **Import from PowerPivot** template as the project type.

Note:

* To create a new tabular model project from a PowerPivot for Excel file

In SQL Server Data Tools, on the File menu, click New, and then click Project.

(box 2)

In the New Project dialog box, under Installed Templates, click Business Intelligence, and then click Import from PowerPivot.

(box 3)

In Name, type a name for the project, then specify a location and solution name, and then click OK.

In the Open dialog box, select the PowerPivot for Excel file that contains the model metadata and data you want to import, and then click Open.

Box 1:

* When creating a new tabular model project by importing from a PowerPivot workbook, the metadata that defines the structure of the workbook is used to create and define the structure of the tabular model project in SQL Server Data Tools. Objects such as tables, columns, measures, and relationships are retained and will appear in the tabular model project as they are in the PowerPivot workbook.

* Analysis Services provides three different approaches for creating a business intelligence semantic model: tabular, multidimensional, and PowerPivot. Tabular solutions use relational modeling constructs such as tables and relationships for modeling data, and the xVelocity in-memory analytics engine for storing and calculating data.

Reference: Import from PowerPivot (SSAS Tabular)

Question: 114

A production SQL Server Analysis Services (SSAS) cube is processed daily. The users query products by using a hierarchy named Products from a dimension named Product.

The DimProduct table in the data source view is used as the source of the Product dimension. The table has the following structure.

```
CREATE TABLE [dbo].[DimProduct](
    [DimensionKey] [int] IDENTITY(1,1) NOT NULL,
    [ProductKey] [int] NOT NULL,
    [ProductName] [varchar](50) NOT NULL,
    [SubCategoryKey] [int] NOT NULL,
    [SubCategoryName] [varchar](50) NOT NULL,
    [CategoryKey] [int] NOT NULL,
    [CategoryName] [varchar](50) NOT NULL
) ON [PRIMARY]
```

The **Product** dimension has three attribute hierarchies:

- **Product**
- **SubCategory**
- **Category**

The Product dimension has three attribute hierarchies:

Product

Subcategory

Category

The attributes have the following relationships defined: Product > Subcategory > Category. Each attribute has a key and a name sourced from the related key and name columns in the DimProduct table.

During processing, you receive the following error message: 'Errors in the OLAP storage engine: A duplicate attribute key has been found when processing: Table: 'dbo_DimProduct', Column: 'SubCategoryKey1', Value: "23'. The attribute is Subcategory'."

You verify that the data is accurate.

You need to ensure that the dimension processes successfully.

What should you do?

- A. Delete the Products hierarchy.
- B. Relate the Subcategory and Category attributes directly to the Product attribute.
- C. Remove the duplicate data from the DimProduct table.
- D. Remove the Subcategory attribute.

Answer: B

Question: 115

You have a SQL Server Analysis Services (SSAS) database named DB1 on a server named Server1. You need to deploy DB1 from Server1 to four other servers. If the destination server already contains a copy of DB1, the database must be updated only. If the destination server does not contain a copy of DB1, the database must be copied to the destination server. What is the best deployment option to use? More than one answer choice may achieve the goal. Select the BEST answer.

- A. Analysis Services Deployment Wizard
- B. Analysis Management Objects (AMO) automation
- C. Backup and Restore
- D. Synchronize Database Wizard

Answer: D

In the Synchronize Database Wizard dialog box, type the name of the source server and source database in the appropriate fields.

Synchronization will occur for source and destination databases that have the same name. If the destination server already has a database that shares the same name as the source database, the destination database will be updated with the metadata and data of the source. If the database does not exist, it will be created on the destination server. Click Next.

Ref: [http://technet.microsoft.com/en-us/library/ms174928\(v=sql.110\).aspx](http://technet.microsoft.com/en-us/library/ms174928(v=sql.110).aspx)

Question: 116

You are deploying an update to a SQL Server Analysis Services (SSAS) cube to a production environment. The production database has been configured with security roles. You need to preserve the existing security roles in the production database. Database roles and their user accounts from the development environment must not be deployed to the production server. Which deployment method should you use?

- A. Use the SQL Server Analysis Services Deployment Wizard.
- B. Backup and restore the database.
- C. Use the SQL Server Analysis Services Migration Wizard.
- D. Deploy the project from SQL Server Data Tools to the production server.

Answer: A

Question: 117

You maintain a multidimensional Business Intelligence Semantic Model (BISM) that was developed with default settings. The model has one cube and the cube has one measure group. The measure group is based on a very large fact table and is partitioned by month. The fact table is incrementally loaded each day with approximately 800,000 new rows. You need to ensure that all rows are available in the cube while minimizing the processing time. Which processing option should you use?

- A. Process Index
- B. Process Data
- C. Process Add
- D. Process Default
- E. Process Clear
- F. Process Full

Answer: C

Question: 118

DRAG DROP

You are developing a SQL Server Analysis Services (SSAS) multidimensional project. The project file includes two cubes named Finance and Operations. The project also includes a dimension named Date. The Date dimension includes two hierarchies named Fiscal and Calendar. The Date dimension has been added to both cubes.

You need to disable the Fiscal hierarchy in the Operations cube without impacting other database objects.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

- | |
|--|
| Open the Date dimension in the dimension designer. |
| Delete the Fiscal hierarchy from the Operations cube. |
| Open the Operations cube in the cube designer. |
| In the Properties window, set the Enabled property to False . |
| In the Dimensions pane of the Cube Structure tab, select the Fiscal hierarchy of the Date dimension. |
| In the Properties window, set the Visible property to False . |
| In the Properties window, set the AttributeHierarchyEnabled property to False . |
| In the Hierarchies pane of the dimension structure tab, select the Fiscal hierarchy. |

Answer:

Box 1:

Open the **Operations** cube in the cube designer.

Box 2:

In the Dimensions pane of the Cube Structure tab, select the **Fiscal** hierarchy of the **Date** dimension.

Box 3:

In the Properties window, set the **AttributeHierarchyEnabled** property to **False**.

Explanation:

Note:

* The value of the AttributeHierarchyEnabled property determines whether an attribute hierarchy is created. If this property is set toFalse, the attribute hierarchy is not created and the attribute cannot be used as a level in a user hierarchy; the attribute hierarchy exists as a member property only. However, a disabled attribute hierarchy can still be used to order the members of another attribute. If the value of the AttributeHierarchyEnabled property is set toTrue, the value of the AttributeHierarchyVisible property determines whether the attribute hierarchy is visible independent of its use in a user-defined hierarchy.

* To set the attribute hierarchy properties in the Employee dimension

Switch to Dimension Designer for the Employee dimension, and then click the Browser tab.

Verify that the following attribute hierarchies appear in the Hierarchy list:

Base Rate
Birth Date
Login ID
Manager SSN
SSN

Switch to the Dimension Structure tab, and then select the following attributes in the Attributes pane, by using the CTRL key to select multiple attributes at the same time:

Base Rate
Birth Date
Login ID
Manager SSN
SSN

In the Properties window, set the value of the AttributeHierarchyEnabled property to False for the selected attributes.
Etc.

Reference: Hiding and Disabling Attribute Hierarchies

Question: 119

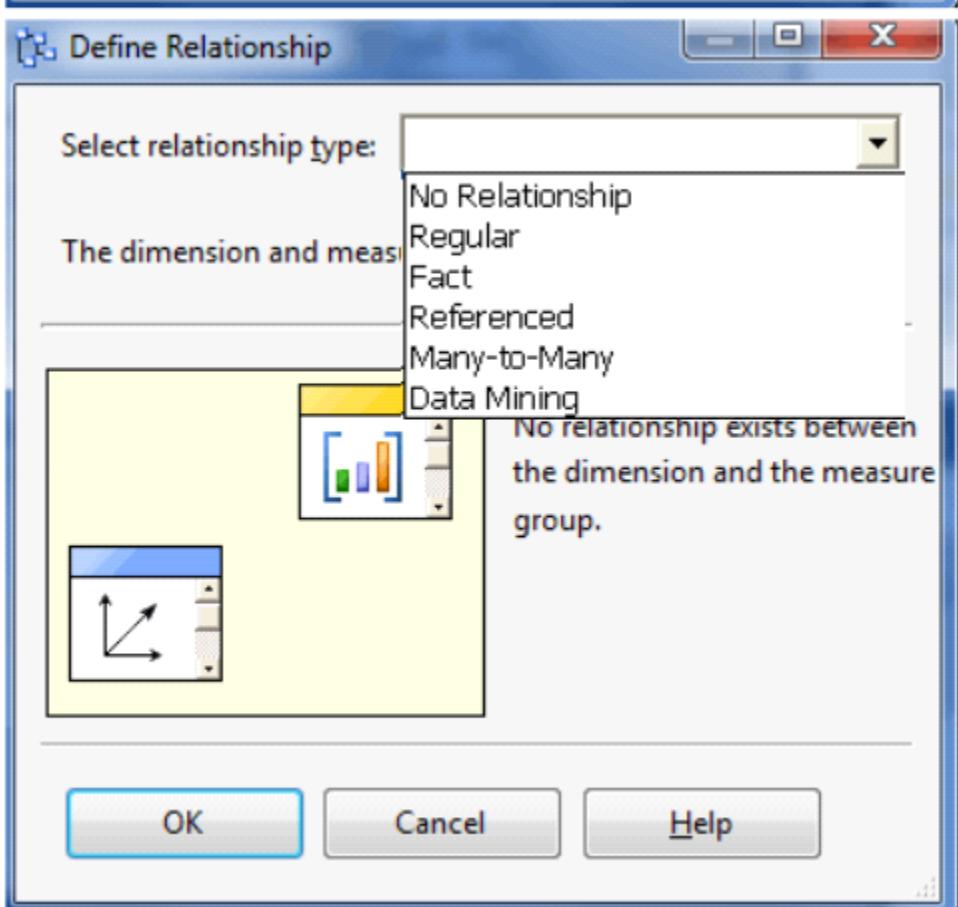
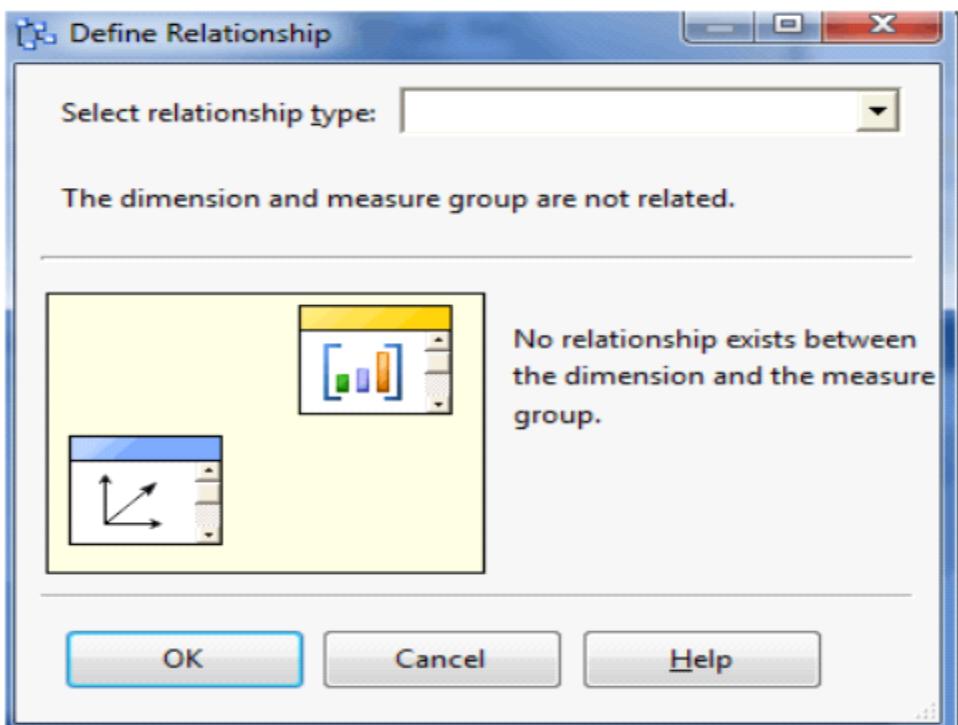
HOTSPOT

You are developing a SQL Server Analysis Services (SSAS) cube.

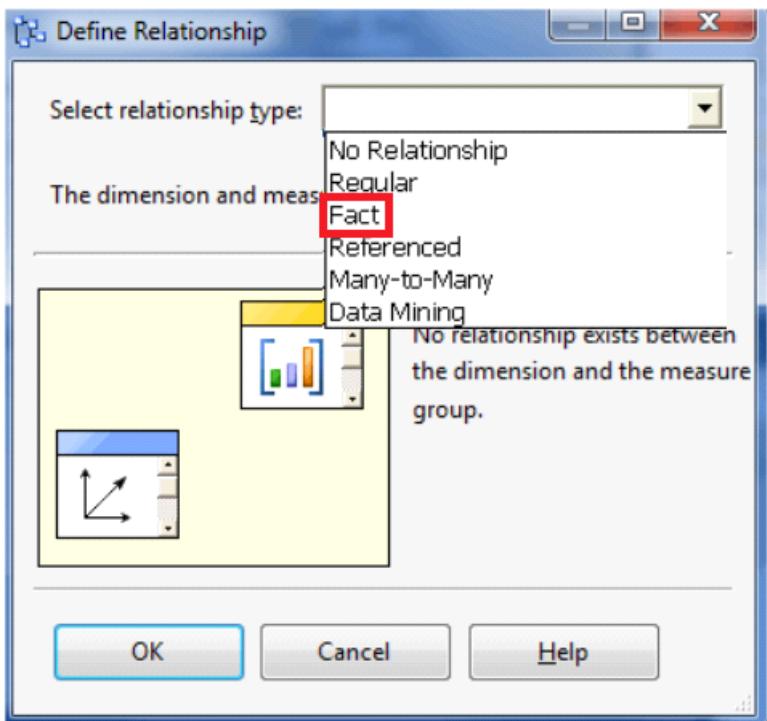
You create a degenerate dimension.

You need to define a relationship type for the dimension.

Which relationship type should you select? (To answer, configure the appropriate option or options in the dialog box in the answer area.)



Answer:

**Question: 120**

You are modifying a SQL Server Analysis Services (SSAS) cube.

Users of the cube report that the precision for the TransactionCost measure is five digits.

You need to ensure that the TransactionCost measure stores values to two digits of precision.

What should you do?

- A. Use the FormatString measure property to format TransactionCost as Currency.
- B. Add a named calculation in the data source view that casts the data source column to two digits of precision. Bind the TransactionCost measure to the new column.
- C. Add a named query in the data source view that casts the data source column to two digits of precision. Bind the TransactionCost measure to the new query.
- D. Use the MeasureExpression measure property to change the precision of TransactionCost to two digits.
- E. Use the FormatString measure property to format TransactionCost as #, ##0.00;-#,##0.00.

Answer: B

Question: 121**HOTSPOT**

You are developing a SQL Server Analysis Services (SSAS) cube.

Revenue must be compared to a goal and described by a status and a trend. Revenue, goal, status, and trend will be defined by Multidimensional Expressions (MDX) expressions.

You need to add the Revenue indicator.

Which tab should you select? (To answer, select the appropriate tab in the work area.)

**Answer:****Question: 122**

You are modifying a SQL Server Analysis Services (SSAS) cube that aggregates mobile phone usage data from a Microsoft Azure SQL Database database. The existing database contains a device dimension.

The Research and Development team has requested that capabilities be added to the database.

The capabilities must meet the following requirements:

- A device member must be able to have multiple capability members.
- A capability member must be able to belong to several device members.
- The Research and Development team must be able to create new capabilities every quarter in the data source.

You need to implement the appropriate solution to meet the requirements while ensuring that the amount of development and maintenance time is minimized.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create a dimension named Capability Name and then configure a many-to-many relationship.
- B. Add an attribute hierarchy for each capability to the customer dimension. Configure each hierarchy to have two members named Yes and No.
- C. Create a dimension named Capability Name and then configure a regular relationship.
- D. Add an attribute hierarchy named Capability Name to the customer dimension.

Answer: A

Question: 123

You are adding dimensions and a new measure group named Service Quality to an existing SQL Server Analysis Services (SSAS) cube.

Date is one of the existing dimensions in the SSAS database. The underlying fact table for the measure group is associated with multiple dates, including FirstServiceDate, SecondServiceDate, and ThirdServiceDate.

You need to ensure that users can slice the Service Quality measures by FirstServiceDate, SecondServiceDate, and ThirdServiceDate. You also need to ensure that the time required to process the database is minimized.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Create cube dimensions named FirstServiceDate, SecondServiceDate, and ThirdServiceDate by using the existing date dimension in the database.
- B. Create three independent dimensions named FirstServiceDate, SecondServiceDate, and ThirdServiceDate and then link all three of them to the Service Quality measure group.
- C. Create three new perspectives named FirstServiceDate, SecondServiceDate, and ThirdServiceDate.
- D. Add three attributes named FirstServiceDate, SecondServiceDate, and ThirdServiceDate to the existing date dimension in the database.

Answer: A

Question: 124

You develop a SQL Server Analysis Services (SSAS) stored procedure.

You need to ensure that developers can create Multidimensional Expressions (MDX) calculations that use the stored procedure.

What should you do?

- A. Register the assembly on the SSAS server instance.
- B. Copy the assembly to the SSAS installation directory and register it by using the CREATE ASSEMBLY T-SQL command.
- C. Register the assembly on the SSAS server by using regedit.exe.
- D. Register the assembly on the SharePoint server by using regedit.exe.

Answer: A

Question: 125

DRAG DROP

You have a cube named Cube1 that contains the sales data for your company.

You plan to build a report based on the cube.

You need to write an MDX expression that returns the total sales from the first month of the 2009 fiscal year and the total sales from the same period of the 2008 fiscal year.

Which code segments should you insert at line 03 and line 05? To answer, drag the appropriate code segments to the correct lines. Each code segments may be used once, more than once, or not at all. You may need to drag the split bar between panes or scroll to view content.

Code segments

```
OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] )
```

```
ParallelPeriod ( [Date].[Fiscal].[Fiscal Year] , -1 ,
                  OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] ) )
```

```
ParallelPeriod ( [Date].[Fiscal].[Fiscal Year] , 1 ,
                  OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] ) )
```

```
OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] )
```

```
ParallelPeriod ( [Date].[Fiscal].[Month] , 1 ,
                  OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] ) )
```

```
ParallelPeriod ( [Date].[Fiscal].[Month] , -1 ,
                  OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] ) )
```

Answer Area

01. `SELECT [Measures].[Sales Amount] ON ROWS`

02. `, {`

03. `code segment`

04. `,`

05. `code segment`

06. `} ON COLUMNS`

07. `FROM [Cube1];`

Answer:

Answer Area

```

01. SELECT [Measures].[Sales Amount] ON ROWS

02. ,{

03.     OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal Year].&[2009] )

04.     ,

05.     ParallelPeriod ( [Date].[Fiscal].[Fiscal Year] , 1 ,
06.                         OpeningPeriod ( [Date].[Fiscal].[Month] , [Date].[Fiscal].[Fiscal
07.                         Year].&[2009] ) )
08.         } ON COLUMNS

09. FROM [Cube1];

```

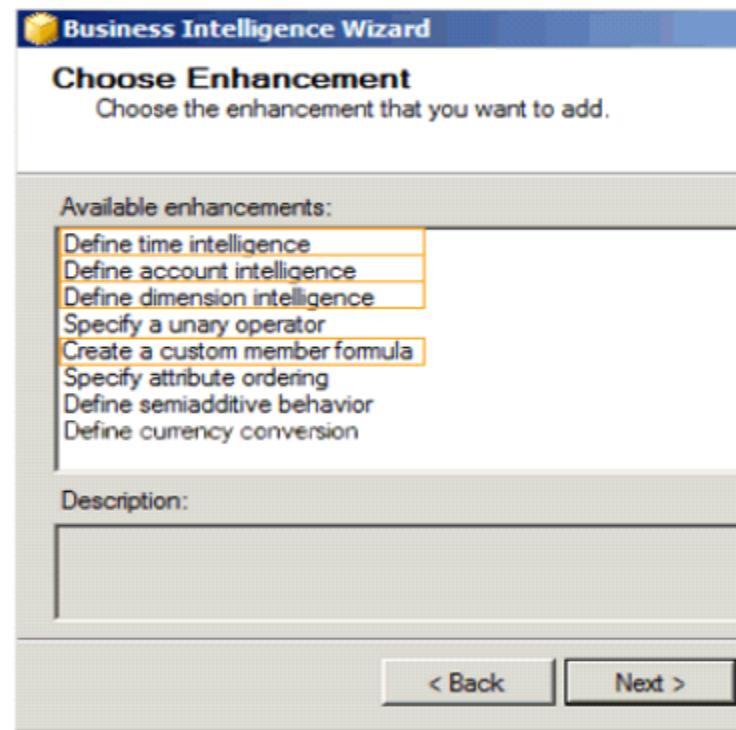
Question: 126**HOTSPOT**

A sales cube contains two years of data.

The sales team must see year-over-year (YOY) and month-over-month (MOM) sales metrics.

You need to modify the cube to support the sales team's requirements.

Which Business Intelligence Wizard enhancements should you use? (To answer, configure the appropriate option or options in the dialog box in the answer area.)

**Answer:**

Available enhancements:

- Define time intelligence**
- Define account intelligence
- Define dimension intelligence
- Specify a unary operator
- Create a custom member formula
- Specify attribute ordering
- Define semiadditive behavior
- Define currency conversion

Description:

< Back Next >

Question: 127

You are developing a SQL Server Analysis Services (SSAS) tabular project. A model defines a measure named Profit and includes a table named Date. The table includes year, semester, quarter, month, and date columns. The Date column is of data type Date. The table contains a set of contiguous dates.

You need to create a measure to report on year-over-year growth of profit.

What should you do? (Each answer presents a complete solution. Choose all that apply.)

- A. Define the following calculation.

```
Year Over Year Profit Growth:=CALCULATE([Profit], DATEADD('Date'[Date], 1, YEAR))
```

- B. Define the following calculation.

```
Year Over Year Profit Growth:=[Profit] - CALCULATE([Profit], PARALLELPERIOD('Date'[Date], -12, MONTH))
```

- C. Define the following calculation.

```
Year Over Year Profit Growth:=[Profit] - CALCULATE([Profit], SAMEPERIODLASTYEAR('Date'[Date]))
```

- D. Use the Business Intelligence Wizard and then use the **Define time intelligence** enhancement.

- Option A
- Option B
- Option C
- Option D

Answer: B, C, D

Question: 128

You are developing a SQL Server Analysis Services (SSAS) tabular project.

You need to grant permission for salespersons to be able to view only the data based on their sales territory.

What should you do?

- A. Create a member and then create a Multidimensional Expressions (MDX) filter.
- B. Create a member and then create a Data Analysis Expressions (DAX) filter.
- C. Create a role and then create a Multidimensional Expressions (MDX) filter.
- D. Create a role and then create a Data Analysis Expressions (DAX) filter.

Answer: D

Question: 129

You deploy a tabular model in DirectQuery mode to a server named Server1. Server1 has SQL Server Analysis Services (SSAS) installed. The model uses a relational data source named DB1.

You need to ensure that Server1 passes the credentials of the current user browsing the model to DB1. What should you do?

- A. Register a Service Principal Name (SPN) for Server1.
- B. From the SQL Server Data Tools, modify the impersonation settings.
- C. From the Web.config file, modify the impersonation settings.
- D. Register a Service Principal Name (SPN) for DB1.

Answer: A

Question: 130

DRAG DROP

You are developing a SQL Server Reporting Services (SSRS) report to display a list of employees. The report will be embedded into a Microsoft SharePoint Server Web Part Page of the company intranet site.

The report consists of a single table. The design of the table is shown in the following diagram.

Employee	Email Address
[EmployeeName]	[EmailAddress]

You need to configure the EmailAddress detail text box to create a new email message. The email message must be addressed to the email address that was clicked by the user.

Which four actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Open the Text Box Properties window of the **EmailAddress** detail text box.

In the box, enter the =Fields! EmailAddress.Value expression.

Select the **Action** page.

In the box, enter the =UrlEncode("mailto:" & Fields! EmailAddress.Value)expression.

In the box, enter the ="mailto:" & Fields! EmailAddress.Value expression.

Select the **General** page.

Select the **Go to URL** option.

Select the **Go to Bookmark** option.

Answer:

Box 1:

Open the Text Box Properties window of the **EmailAddress** detail text box.

Box 2:

Select the **Action** page.

Box 3:

Select the **Go to URL** option.

Box 4:

In the box, enter the ="mailto:" & Fields! EmailAddress.Value expression.

Explanation:

Note:

- * Select Action on the Text Box Properties dialog box to enable hyperlink options for the text box.

- * Go to URL

Choose this option to define a link to a Web page. Type or select the URL of a Web page or an expression that evaluates to the URL of a Web page. Click the Expression (fx) button to change the expression. This expression can include a field that contains a URL.

Question: 131

HOTSPOT

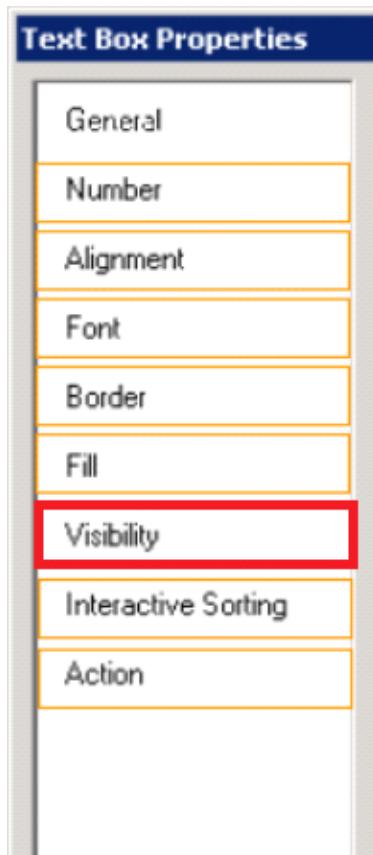
You create a new report in SQL Server Data Tools (SSDT). The report queries a Windows Azure SQL Database database table.

In a report table showing sales by countries and cities, you need to enable users to hide or show cities by clicking the name of a country.

Which property page of the city text box should you use? (To answer, select the appropriate tab in the answer area.)



Answer:



Question: 132

You are developing a SQL Server Reporting Services (SSRS) report. The report includes a dataset with fields named Year, MonthNumber, and RegCount. The report includes a table that displays the number of recorded registration occurrences per year, as shown in the following table.

Year	Reg Count
1995	646,530
1996	728,000
1997	776,255
1998	839,324
1999	867,536
2000	972,305
2001	1,437,642
2002	1,436,908
2003	1,410,109
2004	1,324,998
2005	1,297,629
2006	1,313,613
2007	1,315,046
2008	1,354,499
2009	1,325,585
2010	1,394,674

You need to modify the table to include a graphical item displaying the monthly registration trend to the right of the Reg Count column.

What should you do?

- A. Add an Indicator item to a new column on the right of the Reg Count column. Select the Directional Indicator Type and then assign the MonthNumber field to the Start property.
- B. Add an Indicator item to a new column on the right of the Reg Count column. Select the Directional Indicator Type and then select the MonthNumber field for Value.
- C. Add a Sparkline item to a new column on the right of the Reg Count column. Then select the RegCount field for Values and the MonthNumber field for Series Groups.
- D. Add a text box to a new column on the right of the Reg Count column. Then use a Go to report action to link to a separate report showing the monthly trend.
- E. Add a Sparkline item to a new column on the right of the Reg Count column. Then select the RegCount field for Values and the MonthNumber field for Category Groups.

Answer: E

Question: 133

You manage a SQL Server Reporting Services (SSRS) instance.

An application must pass credentials to the local security authority for Reporting Services.

You need to configure Reporting Services to issue a challenge/response when a connection is made without credentials.

Which authentication type should you configure in the RSReportServer.config file?

- A. RSWindowsKerberos
- B. RSWindowsNegotiate
- C. RSWindowsNTLM
- D. RSWindowsBasic

Answer: D

Ref: <http://msdn.microsoft.com/en-us/library/ms157273.aspx>

Question: 134

You are designing a SQL Server Reporting Services (SSRS) report based on a SQL Server Analysis Services (SSAS) cube.

The cube contains a Key Performance Indicator (KPI) to show if a salesperson's sales are off target slightly off target, or on target.

You need to add a report item that visually displays the KPI status value as a red, yellow, or green circle.

Which report item should you add?

- A. Linear Gauge
- B. Indicator
- C. Data Bar
- D. Radial Gauge
- E. Sparkline

Answer: B

Question: 135

DRAG DROP

You manage a SQL Server Reporting Services (SSRS) instance running in native mode.

You are troubleshooting a performance problem and need to know which reports are frequently executed. You discover that the report server execution logs are empty, despite significant report activity.

You need to ensure that the server is configured for report execution logging.

Which three actions should you perform in sequence? (To answer, move the appropriate actions from the list of actions to the answer area and arrange them in the correct order.)

Enable the **Enable report execution logging** option.

In SQL Server Management Studio (SSMS), connect to the server that runs Report Server.

In the Application Management Group, click **Manage Service Applications**.

Open Reporting Services Configuration Manager.

Open the Server Properties window.

Open SharePoint Central Administration.

Answer:

Box 1:

In SQL Server Management Studio (SSMS), connect to the server that runs Report Server.

Box 2:

Open the Server Properties window.

Box 3:

Enable the **Enable report execution logging** option.

Explanation:

Note: This server is running in NATIVE mode (not Sharepoint mode)

To enable execution logging (in Native mode):

Start SQL Server Management Studio with administrative privileges. For example right-click the Management Studio icon and click 'Run as administrator'.

Connect to the desired report server.

Right-click the server name and click Properties. If the Properties option is disabled, verify you ran SQL Server Management Studio with administrative privileges.

Click the Logging page.

Select Enable report execution Logging.

Ref: <http://msdn.microsoft.com/en-us/library/ms159110.aspx>

Question: 136

You are designing a SQL Server Reporting Services (SSRS) report that sources data from a Microsoft Azure SQL Database database.

The report must display the value and status of a Key Performance Indicator (KPI).

Which report item should you use? (Each answer presents a complete solution. Choose all that apply.)

- A. Indicator
- B. Data Bar
- C. Image
- D. Sparkline
- E. Gauge

Answer: A, E

Question: 137

You are developing a SQL Server Reporting Services (SSRS) sales summary report.

The report header consists of several images.

You need to ensure that the header of the report is hidden when a user exports the report to PDF format.

Which Hidden property expression should you use for the report header? {More than one answer choice may achieve the goal. Select the BEST answer.)

- A. =(Globals!RenderFormat.IsInteractive = False)
- B. =(Globals!RenderFormat.Name = "PDF")
- C. =False
- D. =True

Answer: B

Question: 138

You are developing a new SQL Server Reporting Services (SSRS) report in SQL Server Data Tools (SSDT).

The report must define a report parameter to prompt the user for the business unit. Each business unit has a unique color scheme combination of foreground and background colors.

You need to ensure that all of the text boxes in the table headers use the correct business unit colors.

What should you do? (More than one answer choice may achieve the goal. Select the BEST answer.)

- A. Add one report variable for Color. Assign it with an expression to return the appropriate colors. For each header text box, set the Color and BackgroundColor properties by using the variable.
- B. Add two report variables named Color and BackgroundColor. Assign them with expressions to return the appropriate colors. For each header text box, use expressions to set the Color and BackgroundColor properties by using the variables.
- C. For each header text box, assign expressions to the Color and BackgroundColor properties.
- D. Add two Microsoft Visual C# functions to the code block of the report to implement Color and BackgroundColor functions. For each header text box, use expressions to set the Color and BackgroundColor properties by using the functions.

Answer: D

Question: 140

You are designing a SQL Server Reporting Services (SSRS) report to display product names and their year-to-date (YTD) sales quantity YTD sales quantity values are classified in three bands:

High Sales, Medium Sales, and Low Sales.

You add a table to the report. Then you define two columns based on the fields named ProductName and YTDSalesQuantity.

You need to set the color of the product text to red, yellow, or blue, depending on the value of the YTD sales quantity values.

What should you do?

- A. Use an expression for the TextDecoration property of the text box.
- B. Use an expression for the Style property of the text box.
- C. Add an indicator to the table.
- D. Use an expression for the Font property of the text box.
- E. Use an expression for the Color property of the text box.

Answer: E

Question: 141

A production SQL Server Analysis Services (SSAS) cube is processed daily. The users query facts by using a hierarchy named Geography from a dimension named Geography.

The DimGeography table in the data source view is used as the source of the Geography dimension. The table has the following structure.

```
CREATE TABLE [dbo].[DimGeography] (
    [DimensionKey] [int] IDENTITY (1, 1) NOT NULL,
    [CityKey] [int] NOT NULL,
    [CityName] [varchar] (50) NOT NULL,
    [StateProvinceKey] [int] NOT NULL,
    [StateProvinceName] [varchar] (50) NOT NULL,
    [CountryKey] [int] NOT NULL,
    [CountryName] [varchar] (50) NOT NULL
) ON [PRIMARY]
```

The Geography dimension has three attribute hierarchies:

- City
- State-Province
- Country

The attributes have the following relationships defined: City > State-Province > Country. Each attribute has a key and a name sourced from the related key and name columns in the DimGeography table.

During processing, you receive the following error message:

'Errors in the OLAP storage engine:

A duplicate attribute key has been found when processing:

Table: 'dbo_DimGeography',

Column: 'StateProvinceKey', Value: '23'.

The attribute is State-Province'."

You verify that the data is accurate.

You need to ensure that the dimension Processes successfully.

What should you do?

- A. Delete the Geography hierarchy.
- B. Relate the State-Province and Country attributes directly to the City attribute.
- C. Remove the duplicate data from the DimGeography table.
- D. Remove the State-Province attribute.

Answer: B

Question: 144

You execute the following code:

```
CREATE TABLE dbo.Customers
(
    id int PRIMARY KEY
    CustomerName char (10)
)
```

You create a nonclustered index named IX_CustomerName on the CustomerName column.

You execute the following query:

```
SELECT * FROM dbo.Customers
WHERE LEFT (CustomerName, 1) = 'a'
```

You need to reduce the amount of time it takes to execute the query.

What should you do?

- A. Replace LEFT (CustomerName, 1) = 'a' with CustomerName LIKE 'a%'.
- B. Partition the table and use the CustomerName column for the partition scheme.
- C. Replace LEFT (CustomerName, 1) = 'a' with SUBSTRING (CustomerName, 1/1) = 'a'.
- D. Replace IX_CustomerName with a clustered index.

Answer: A

Explanation:

LEFT() is a function, and it would result in an index scan where like '%' would result in index seek which is more efficient.

References: <http://www.mssqltips.com/sqlservertip/1236/avoid-sql-server-functions-in-the-where-clause-forperformance/>

Question: 145

You are developing a SQL Server Analysis Services (SSAS) multidimensional database. The underlying data source does not have a time dimension table.

You need to implement a time dimension.

What should you do?

- A. Use the SQL Server Data Tools Dimension Wizard and generate a time table in the data source.
- B. Use the SQL Server Data Tools Dimension Wizard and generate a time dimension by using the Use an existing table option.
- C. Create a CSV file with time data and use the DMX IMPORT statement to import data from the CSV file.
- D. Create a script by using a sample time dimension from a different multidimensional database. Then create a new dimension in an existing multidimensional database by executing the script.

Answer: A

In Microsoft SQL Server Analysis Services, you can use the Dimension Wizard in SQL Server Data Tools (SSDT) to create a time dimension when no time table is available in the source database.

Generate a time table in the data source Select this option when you have permission to create objects in the underlying data source. The wizard will then generate a time table and store this table in the data source. The wizard then creates the time dimension from this time table.

Generate a time table on the server Select this option when you do not have permission to create objects in the underlying data source. The wizard will then generate and store a table on the server instead of in the data source. (The dimension created from a time table on the server is called a server time dimension.) The wizard then creates the server time dimension from this table.

References: <https://docs.microsoft.com/en-us/sql/analysis-services/multidimensional-models/create-a-timedimension-by-generating-a-time-table>

Question: 146

You have an application that uses a view to access data from multiple tables.

You need to ensure that you can insert rows into the underlying tables by using the view.

What should you do?

- A. Define the view by using the SCHEMABINDING option.
- B. Define the view by using the CHECK option.
- C. Create an INSTEAD OF trigger on the view.
- D. Materialize the view.

Answer: C

Explanation:

Using INSTEAD OF triggers you can override an INSERT, UPDATE, or DELETE operation on a view. For example, you might define an INSTEAD OF INSERT trigger on a view to replace the standard INSERT statement.

References: <https://msdn.microsoft.com/en-us/library/def01zh2.aspx>

Question: 147

You have a server that has SQL Server 2012 installed.

You need to identify which parallel execution plans are running in serial.

Which tool should you use?

- A. Data Profile Viewer
- B. Database Engine Tuning Advisor
- C. Performance Monitor
- D. Extended Events

Answer: D

Question: 148

You have a database named database1. Database developers report that there are many deadlocks. You need to implement a solution to monitor the deadlocks.

The solution must meet the following requirements:

- Support real-time monitoring.
- Be enabled and disabled easily.
- Support querying of the monitored data.

What should you implement? More than one answer choice may achieve the goal. Select the BEST answer.

- A. An Extended Events session
- B. A SQL Server Profiler template
- C. Log errors by using trace flag 1204
- D. Log errors by using trace flag 1222

Answer: A

There are a few ways you can track down queries that are causing deadlocks. For example, you can use the Deadlock Graph as shown in the previous tip SQL Server Profiler Graphical Deadlock Chain. Another solution is using a trace flag to write the deadlock information to the error log. You can also implement trace flag 1222 to do just that.

References: <https://www.mssqltips.com/sqlservertip/2130/finding-sql-server-deadlocks-using-trace-flag-1222/>

Question: 149

You are developing a multidimensional project that includes a dimension named Organization.

The dimension is based on the DimOrganization table in the data warehouse.

The following diagram illustrates the table design.

DimOrganization	
PK	OrganizationKey
FK2	ParentOrganizationKey PercentageOfOwnership OrganizationName ParentOrganizationName
FK1	CurrencyKey

The Organization dimension includes a parent-child hierarchy named Organizations.

The dimension includes the following dimension attributes:

Organization, which is a key attribute

Organizations, which defines the parent-child hierarchy

Currency Code, which is a regular attribute

When users browse the dimension, three hierarchies are visible to them.

You need to ensure that the Organization hierarchy is not visible to users.

What should you do?

- A. Set the AttributeHierarchyDisplayFolder property to Null for the Organization attribute.
- B. Delete the Organization attribute.
- C. Set the AttributeHierarchyEnabled property to False for the Organization attribute.
- D. Set the AttributeHierarchyVisible property to False for the Organization attribute.

Answer: D

Explanation:

The DimensionAttribute.AttributeHierarchyVisible Property gets or sets whether the attribute hierarchy is visible to client applications.

References: <https://msdn.microsoft.com/en-us/library/microsoft.analysisservices.dimensionattribute.attributehierarchyvisible.aspx>

Question: 150

You are modifying a SQL Server Analysis Services (SSAS) multidimensional database.

You have identified a dimension that is no longer used by any cubes.

You need to delete the dimension.

What should you do?

- A. Write a Multidimensional Expressions (MDX) command to drop the dimension from the database.
- B. Write a Data Mining Extensions (DMX) command to drop the dimension from the database.
- C. Write a T-SQL command to drop the dimension from the database.
- D. Delete the dimension by using SQL Server Management Studio Object Explorer.

Answer: D

To delete a dimension in SQL Server Data Tools.

In Solution Explorer, right-click the dimension that you want to delete, and then click Delete.

References: <https://docs.microsoft.com/en-us/sql/analysis-services/multidimensional-models/databasedimensions-modify-or-delete-a-database-dimension-in-solution-explorer>

Question: 151

You have a database for a mission-critical web application. The database is stored on a SQL Server 2012 instance and is the only database on the instance. The application generates all T-SQL statements dynamically and does not use stored procedures.

You need to maximize the amount of memory available for data caching.

Which advanced server option should you modify?

- A. Scan for Startup Procs
- B. Allow Triggers to Fire Others
- C. Enable Contained Databases
- D. Optimize for Ad hoc Workloads

Answer: C

Explanation:

A contained database is a database that is isolated from other databases and from the instance of SQL Server that hosts the database. SQL Server helps user to isolate their database from the instance in 4 ways.

Much of the metadata that describes a database is maintained in the database. (In addition to, or instead of, maintaining metadata in the master database.)

All metadata are defined using the same collation.

User authentication can be performed by the database, reducing the databases dependency on the logins of the instance of SQL Server.

The SQL Server environment (DMV's, XEvents, etc.) reports and can act upon containment information.

References: <https://docs.microsoft.com/en-us/sql/relational-databases/databases/contained-databases>

Question: 152

You plan to design an application that temporarily stores data in a SQL Azure database.

You need to identify which types of database objects can be used to store data for the application. The solution must ensure that the application can make changes to the schema of a temporary object during a session.

Which type of objects should you identify?

- A. Common table expressions (CTEs)
- B. Table variables
- C. Temporary tables
- D. Temporary stored procedures

Answer: C

Explanation:

There are two types of temporary tables: local and global. Local temporary tables are visible only to their creators during the same connection to an instance of SQL Server as when the tables were first created or referenced. Local temporary tables are deleted after the user disconnects from the instance of SQL Server. Global temporary tables are visible to any user and any connection after they are created, and are deleted when all users that are referencing the table disconnect from the instance of SQL Server.

References: [https://technet.microsoft.com/en-us/library/ms186986\(v=sql.105\).aspx](https://technet.microsoft.com/en-us/library/ms186986(v=sql.105).aspx)

Question: 153

You have a SQL Server 2012 instance that hosts a single-user database. The database does not contain user-created stored procedures or user-created functions. You need to minimize the amount of memory used for query plan caching.

Which advanced server option should you modify?

- A. Enable Contained Databases
- B. Allow Triggers to Fire Others
- C. Optimize for Ad hoc Workloads
- D. Scan for Startup Procs

Answer: C

Explanation:

The optimize for ad hoc workloads option is used to improve the efficiency of the plan cache for workloads that contain many single use ad hoc batches. When this option is set to 1, the Database Engine stores a small compiled plan stub in the plan cache when a batch is compiled for the first time, instead of the full compiled plan. This helps to relieve memory pressure by not allowing the plan cache to become filled with compiled plans that are not reused.

References: <https://docs.microsoft.com/en-us/sql/database-engine/configure-windows/optimize-for-ad-hocworkloads-server-configuration-option>

Question: 154

You have a SQL Server 2012 database named DB1. You have a backup device named Device1.

You discover that the log file for the database is full. You need to ensure that DB1 can complete transactions.

The solution must not affect the chain of log sequence numbers (LSNs).

Which code segment should you execute?

- A. BACKUP LCG DB1 TO Device1 WITH COPY_ONLY
- B. BACKUP LOG DB1 TO Device1
- C. BACKUP LOG DB1 TO Device1 WITH NCRECCVERY
- D. BACKUP LOG D31 TO Device1 WITH TRUNCATE ONLY

Answer: B

Explanation:

The transaction log LSN chain is not affected by a full or differential database backup.

References: <https://www.mssqltips.com/sqlservertip/3209/understanding-sql-server-log-sequence-numbers-for-backups/>